

## Before we start...

Access the RStudio Cloud Project (Link in chat)

- Workshop - Introduction to R & Data Wrangling

- 
1. Download R
  2. Download RStudio
  3. Download the `titanic` dataset

<https://web.stanford.edu/class/archive/cs/cs109/cs109.1166/problem12.html>

# Introduction to R and Data Wrangling (with tidyverse)



# Workshop Housekeeping



## Questions?

Use the Q&A, chat, or raise hand feature.



## Feedback Survey

After the session, please fill out the Qualtrics survey (link in chat).



## Credits

Many images are sourced from the teaching team at Harvard Chan Bioinformatics Core (HBC).

Content was similarly inspired by HBC.

Slides with `(HBC Source)` in the bottom corner indicate the image/table source.

Original source: <https://hbctraining.github.io/Training-modules/IntroR/>

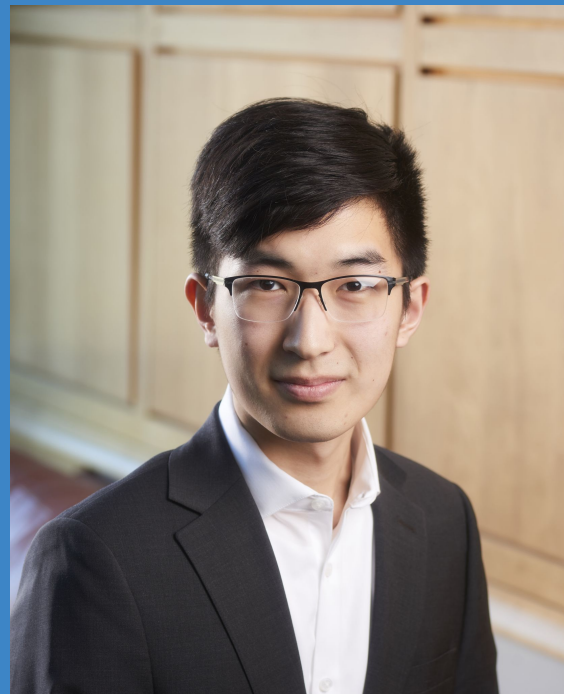
# About Me

- Master's in Applied Statistics
- Bachelor's in Computational Statistics

## R Experience:

- Self-taught for research (2017-Present)
- Statistics Courses (STAT 184/380)
- Other Projects

Research Consultant, University Libraries



David Chen

[dzc89@psu.edu](mailto:dzc89@psu.edu)

Research Consultant  
[github.com/TheDavidChen](https://github.com/TheDavidChen)

# Introduce Yourself!

Name, Major, Year, Favorite Letter



# Agenda

## Introduction to R

- Why R?
- The RStudio Interface
- Running Code
- Variable Assignment
- Data Types
- Data Structures
- Importing Datasets



# Agenda

## Introduction to R

- Why R?
- The RStudio Interface
- Running Code
- Variable Assignment
- Data Types
- Data Structures
- Importing Datasets

## Data Wrangling with Tidyverse

- Why Tidyverse?
- The Fundamental Functions
  - Subset data
  - Create new variables
  - Grouped summaries
  - Combining datasets
- Exporting Datasets



## Why R?

- Free/Open Source
- Platform Agnostic
- Designed for Data Analysis
- Customizable Data Visualizations
- Reproducibility
- Big Data!



Logo: The R Foundation

## Motivation - Max Number of Rows

**Excel: 1,048,576**

**R:  $2^{48} =$**

**281,474,976,710,656**

# RStudio

- Free/Open Source
- Platform agnostic
- Interface for R

Reasons to use  
Base R instead:





A Comment on Time

# Pace vs Content tradeoff



## General Comments

Expect a learning curve. Expect to struggle.

Errors are the best for learning. If they happen:

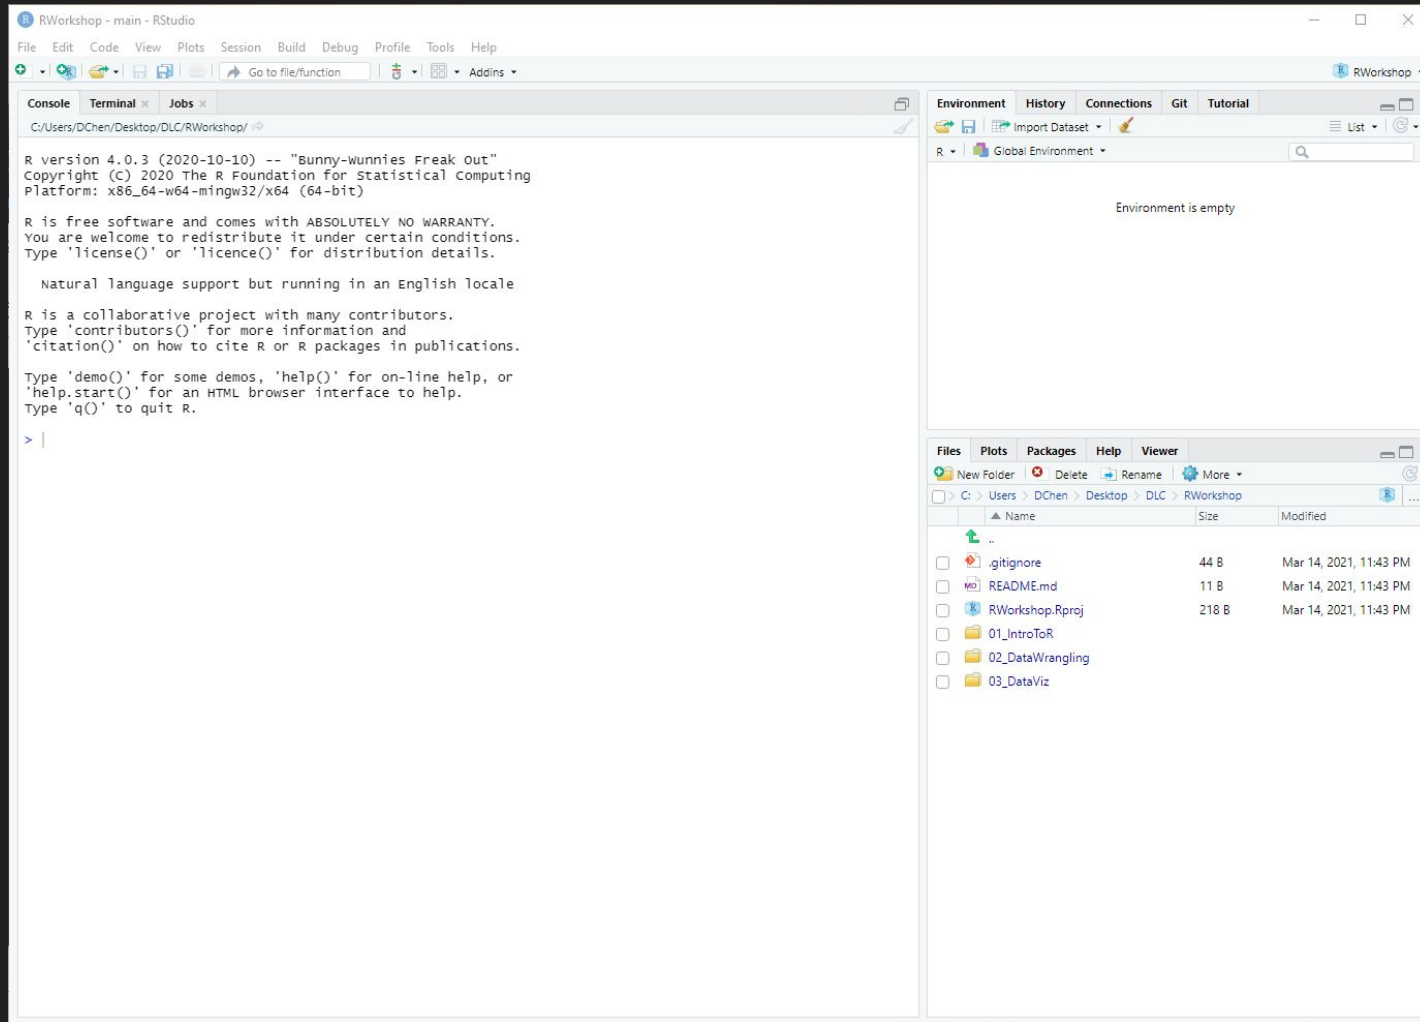
- Read the error and try to solve it
- Share the code that gave you an error
- Share the error message



RStudio is a trademark of RStudio, PBC

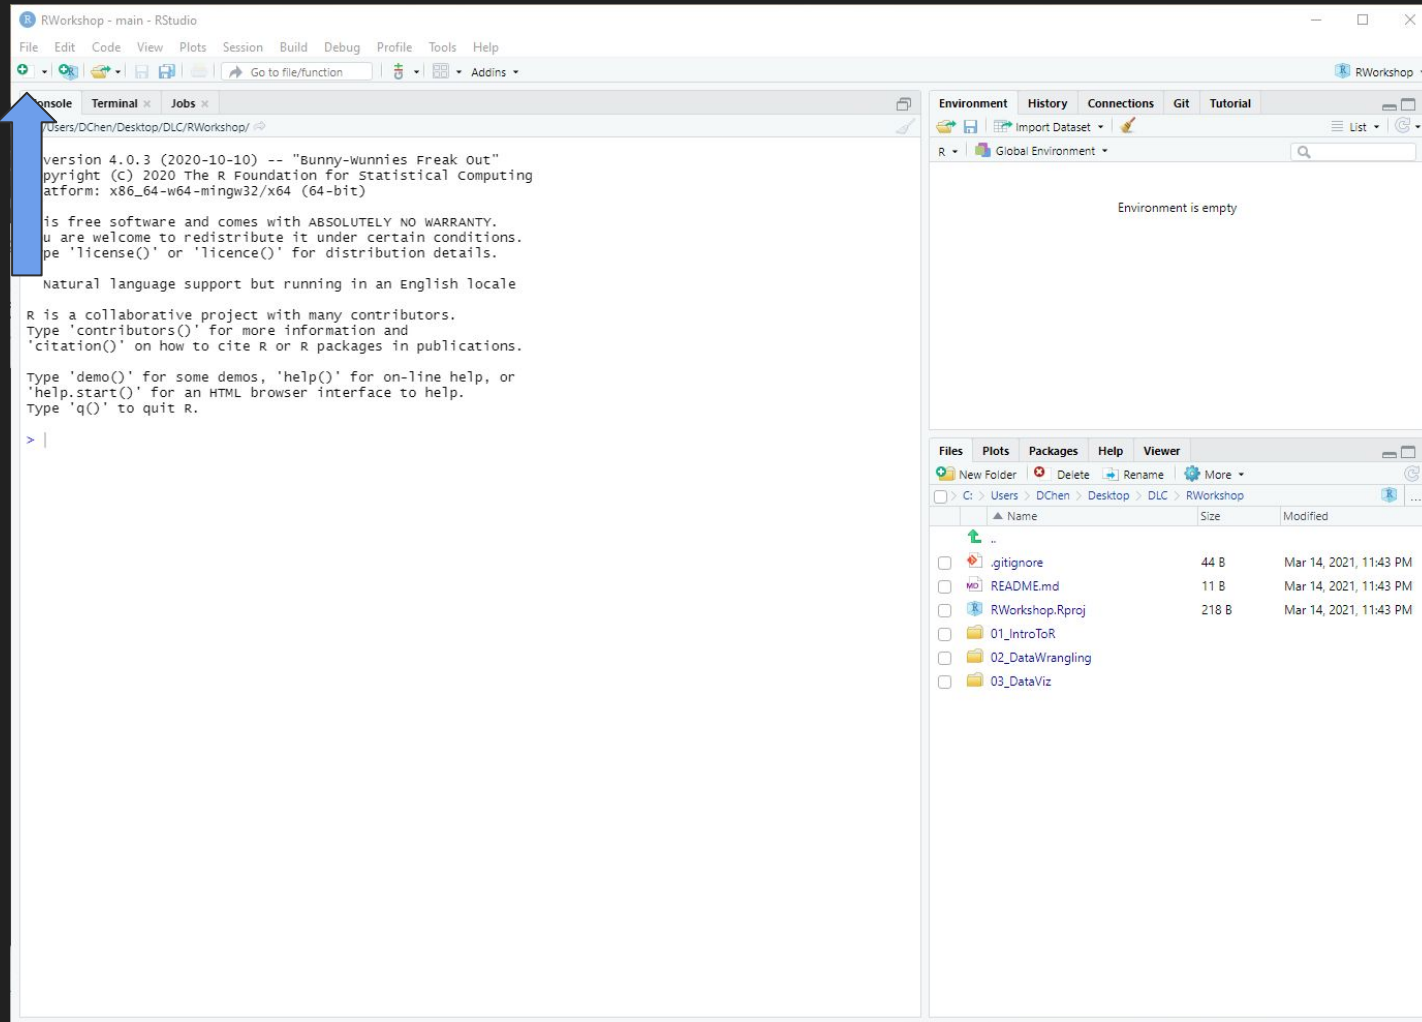
# RStudio

# Open RStudio!



# RStudio

# Open RStudio!



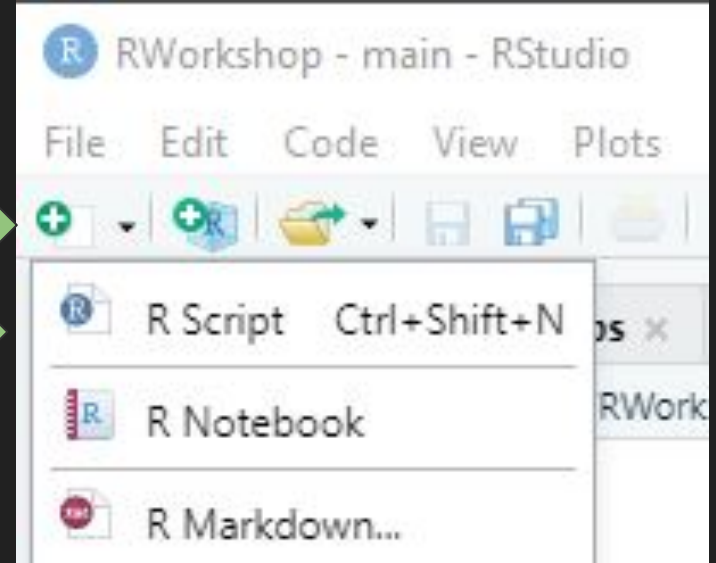


# RStudio

## Create an R Script

Three options:

- Click the top left button
- Click `File` -> New File -> R Script`
- Press `Ctrl`+`Shift`+`N`



# RStudio

The screenshot displays the RStudio IDE interface. The main editor window is titled 'Untitled1' and is currently empty. The top menu bar includes File, Edit, Code, View, Plots, Session, Build, Debug, Profile, Tools, and Help. The top toolbar contains icons for opening files, saving, running, and other standard IDE functions. The right sidebar is divided into two sections. The top section, labeled 'Environment', shows the 'Global Environment' and indicates that the environment is empty. The bottom section, labeled 'Files', shows the file explorer for the project directory 'C:\Users\DChe\Desktop\DL\RWorkshop'. It lists several files and folders, including '.gitignore', 'README.md', 'RWorkshop.Rproj', '01\_IntroToR', '02\_DataWrangling', and '03\_DataViz'. The bottom panel is the 'Console', which shows the R version 4.0.3 (2020-10-10) and the R Foundation for Statistical Computing copyright notice. It also displays the R startup message, including the license information and the natural language support but running in an English locale. The console prompt is '> |'.

R version 4.0.3 (2020-10-10) -- "Bunny-wunnies Freak Out"  
Copyright (C) 2020 The R Foundation for Statistical Computing  
Platform: x86\_64-w64-mingw32/x64 (64-bit)

R is free software and comes with ABSOLUTELY NO WARRANTY.  
You are welcome to redistribute it under certain conditions.  
Type 'license()' or 'licence()' for distribution details.

Natural language support but running in an English locale

R is a collaborative project with many contributors.  
Type 'contributors()' for more information and  
'citation()' on how to cite R or R packages in publications.

Type 'demo()' for some demos, 'help()' for on-line help, or  
'help.start()' for an HTML browser interface to help.  
Type 'q()' to quit R.

> |

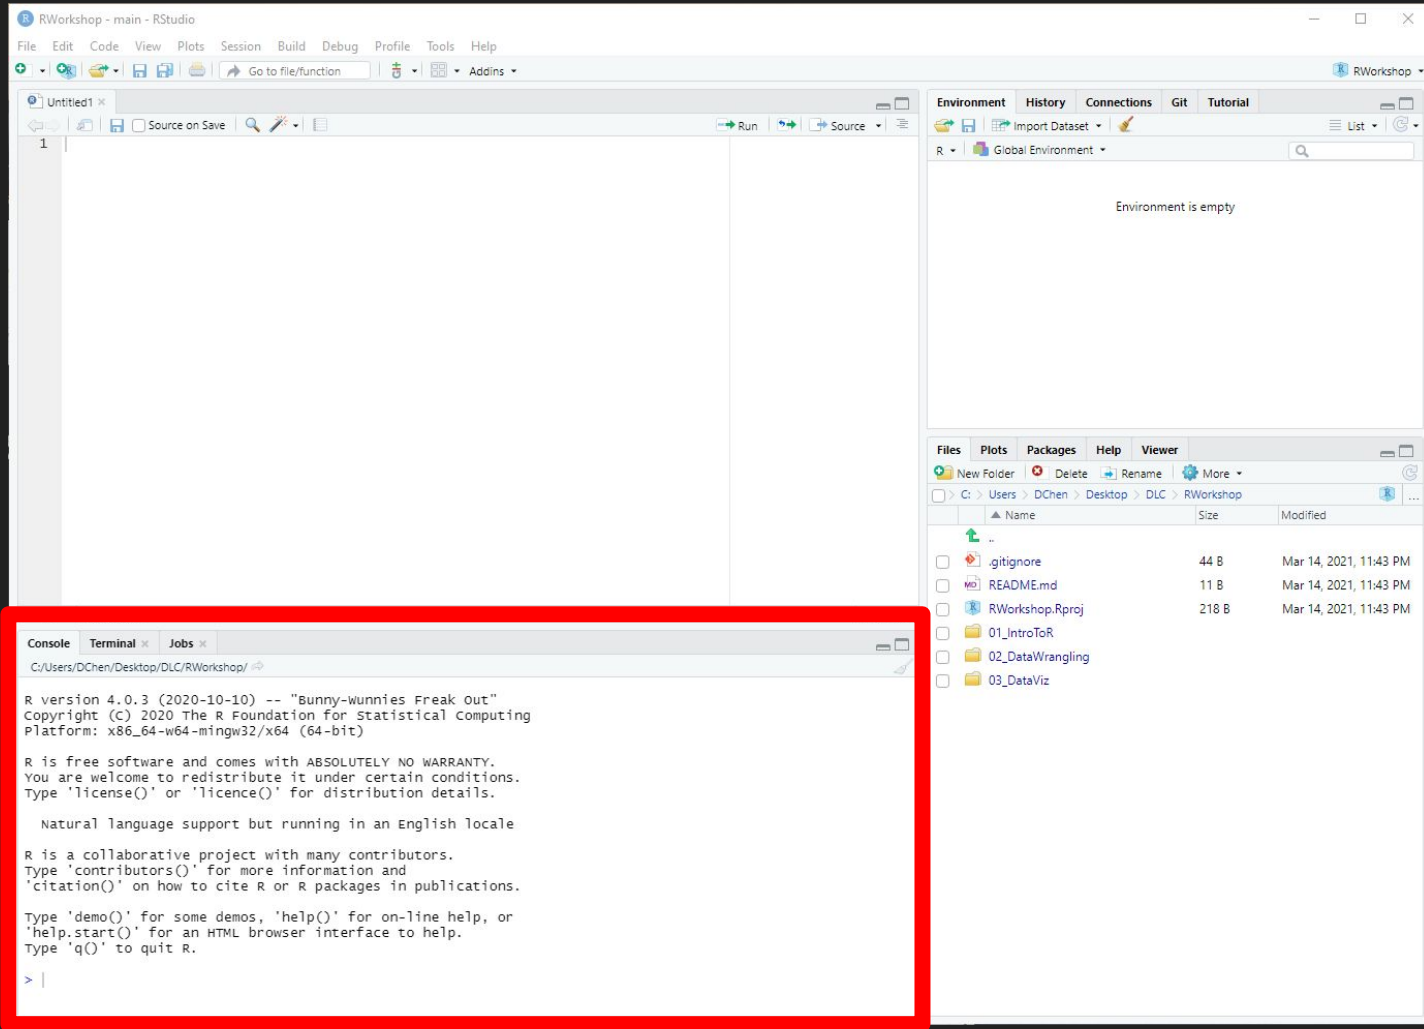
# The Console

Run Code

Show Output

Show Errors

Code Unsaved

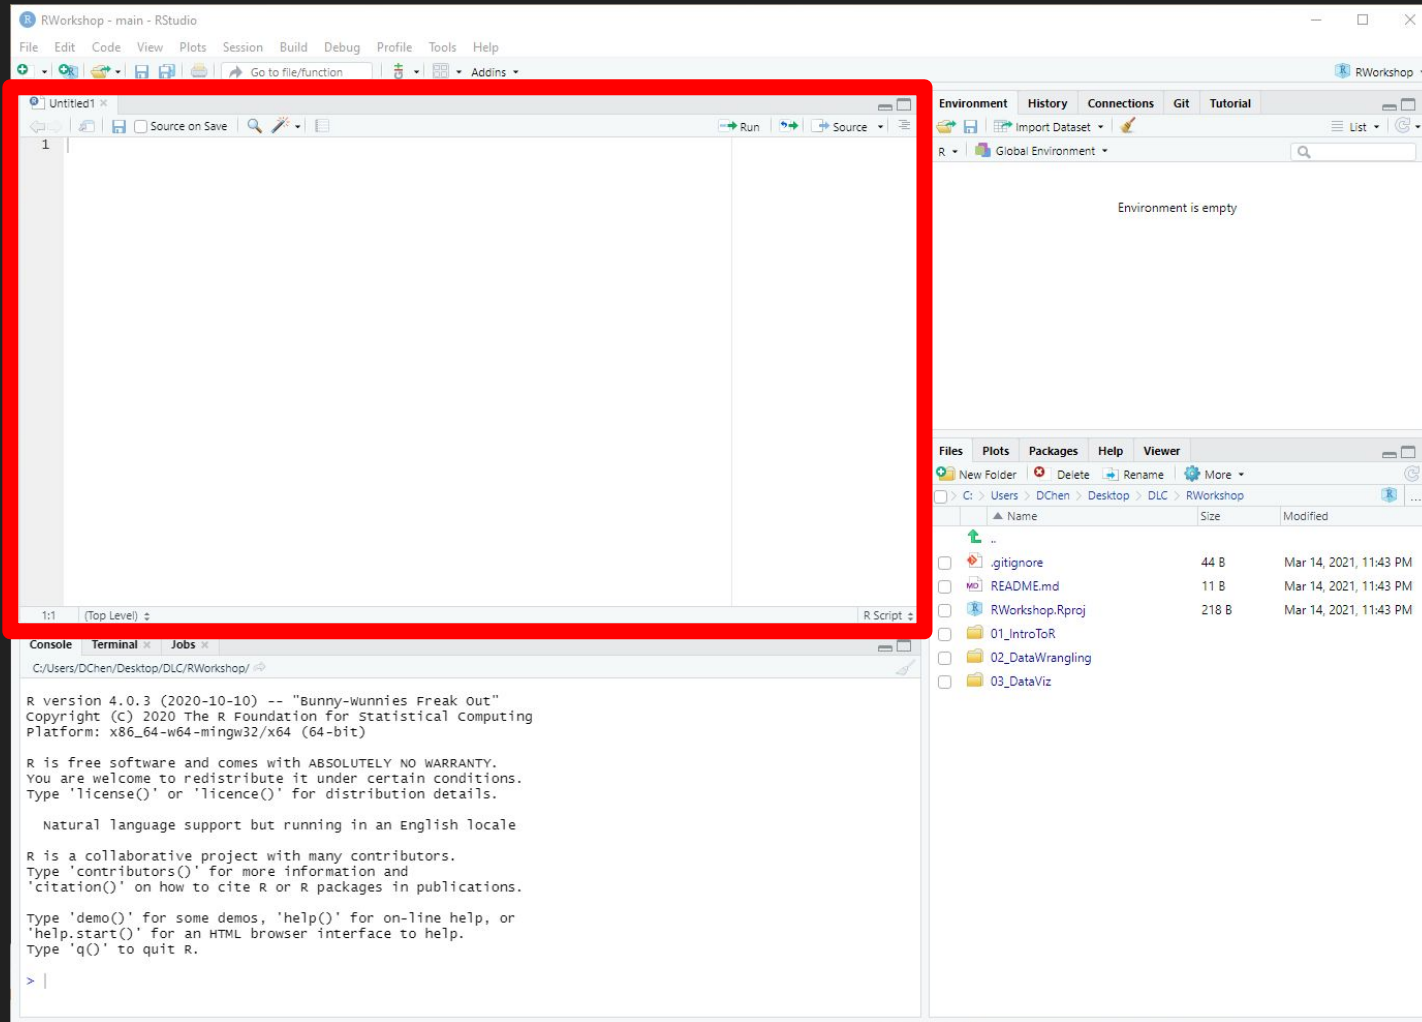


# Script Editor

Type code

Run code (output  
in console)

Saves all code



# Environment

## Import Datasets

Shows all  
datasets and  
defined variables

View imported  
datasets

The screenshot displays the RStudio interface. The main window shows an empty R script file named 'Untitled1'. The bottom pane is split into a Console and a Terminal. The Console shows the R version 4.0.3 (2020-10-10) and copyright information. The Terminal shows the R startup message. The right sidebar contains two panes: the Environment pane and a file explorer. The Environment pane is highlighted with a red border and shows 'R' and 'Global Environment' with the message 'Environment is empty'. The file explorer shows the directory 'C:\Users\DChen\Desktop\DL\ RWorkshop' with files like '.gitignore', 'README.md', 'RWorkshop.Rproj', '01\_IntroToR', '02\_DataWrangling', and '03\_DataViz'.

RStudio - main - RStudio

File Edit Code View Plots Session Build Debug Profile Tools Help

Go to file/function

Imports

Environment History Connections Git Tutorial

R Global Environment

Environment is empty

1:1 (Top Level) R Script

Console Terminal Jobs

C:\Users\DChen\Desktop\DL\RWorkshop\

R version 4.0.3 (2020-10-10) -- "Bunny-wunnies Freak Out"  
Copyright (C) 2020 The R Foundation for Statistical Computing  
Platform: x86\_64-w64-mingw32/x64 (64-bit)

R is free software and comes with ABSOLUTELY NO WARRANTY.  
You are welcome to redistribute it under certain conditions.  
Type 'license()' or 'licence()' for distribution details.

Natural language support but running in an English locale

R is a collaborative project with many contributors.  
Type 'contributors()' for more information and  
'citation()' on how to cite R or R packages in publications.

Type 'demo()' for some demos, 'help()' for on-line help, or  
'help.start()' for an HTML browser interface to help.  
Type 'q()' to quit R.

> |

Files Plots Packages Help View

New Folder Delete Rename More

C:\Users\DChen\Desktop\DL\ RWorkshop

Name	Size	Modified
..		
.gitignore	44 B	Mar 14, 2021, 11:43 PM
README.md	11 B	Mar 14, 2021, 11:43 PM
RWorkshop.Rproj	218 B	Mar 14, 2021, 11:43 PM
01_IntroToR		
02_DataWrangling		
03_DataViz		

# Environment

Import Datasets

Shows all  
datasets and  
defined variables

View imported  
datasets

The screenshot shows the RStudio interface with the Environment pane open. The Environment pane displays the Global Environment with a search bar. Below the search bar, the 'Data' section shows the 'Titanic' dataset with 891 observations and 12 variables. A table of values is shown below, with 'x' having a value of 5. The Files pane on the right shows the project structure with folders for '01\_IntroToR', '02\_DataWrangling', and '03\_DataViz'. The Console pane at the bottom shows the current directory path: C:/Users/DChen/Desktop/DLC/RWorkshop/.

RStudio - main - RStudio

File Edit Code View Plots Session Build Debug Profile Tools Help

Go to file/function

Run Source

Environment History Connections Git Tutorial

R Global Environment

Data

Titanic 891 obs. of 12 variables

Values

x 5

Files Plots Packages Help Viewer

New Folder Delete Rename More

C: > Users > DChen > Desktop > DLC > RWorkshop

Name	Size	Modified
..		
.gitignore	44 B	Mar 14, 2021, 11:43 PM
README.md	11 B	Mar 14, 2021, 11:43 PM
RWorkshop.Rproj	218 B	Mar 14, 2021, 11:43 PM
01_IntroToR		
02_DataWrangling		
03_DataViz		

Console Terminal Jobs

C:/Users/DChen/Desktop/DLC/RWorkshop/

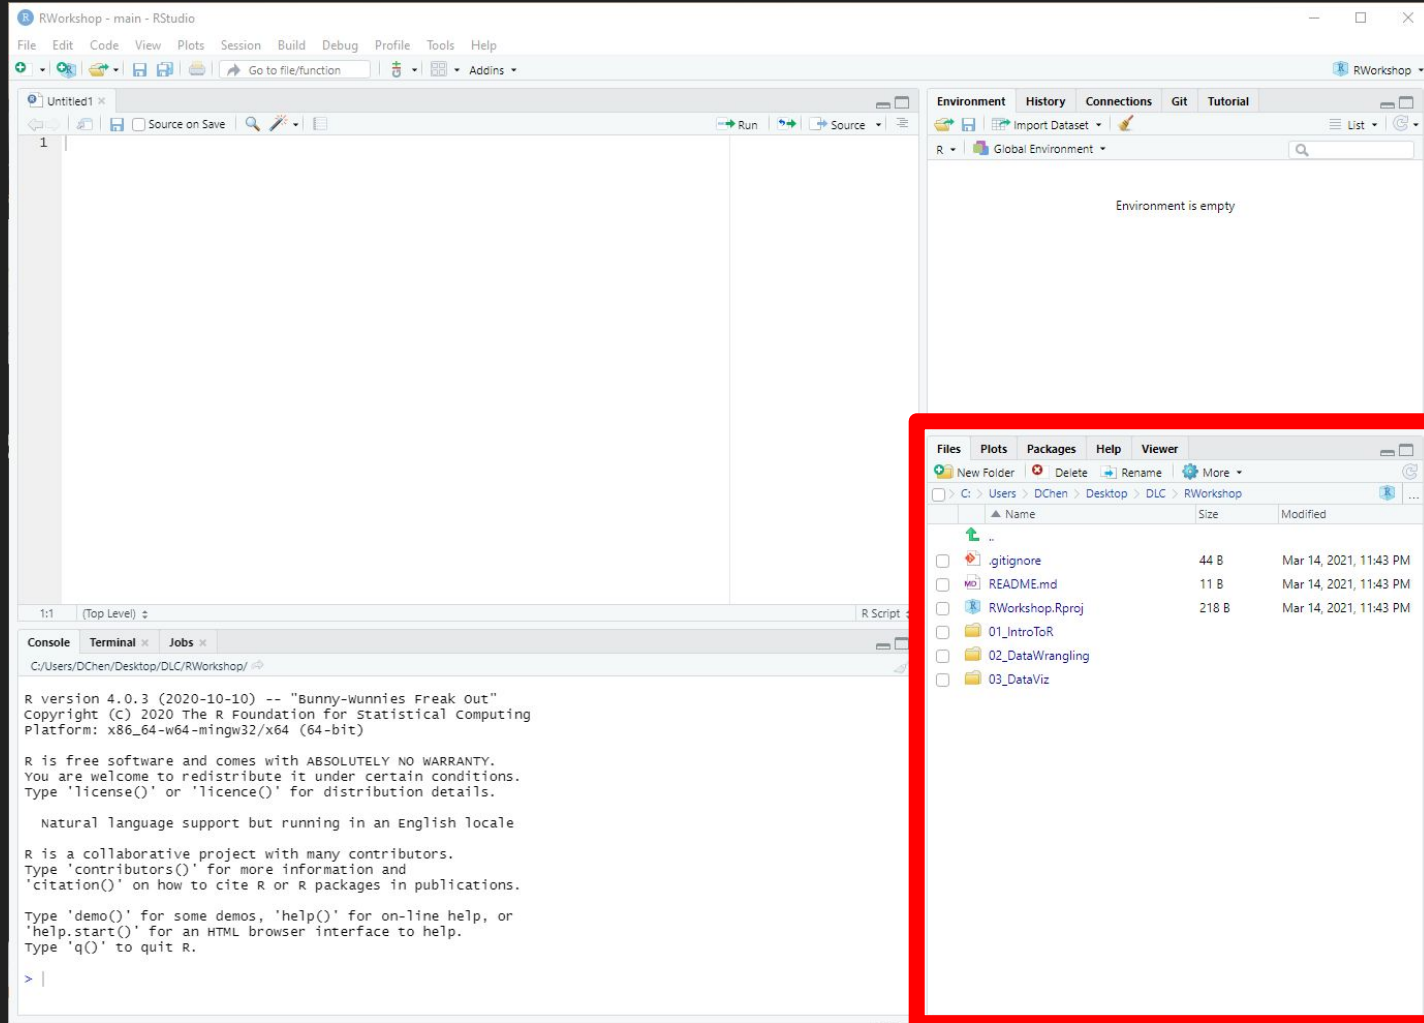
>

# Misc.

## Access Files

## View created plots

## Read help documentation



The screenshot displays the RStudio environment. The main editor window shows a blank R script. The console at the bottom displays the R version 4.0.3 startup message and the R help text. A file explorer window is open in the bottom right corner, showing the contents of the RStudio project directory. The file explorer window is highlighted with a red border.

**RStudio Interface:**

- Menu Bar:** File, Edit, Code, View, Plots, Session, Build, Debug, Profile, Tools, Help.
- Toolbar:** Includes icons for running code, saving, and other standard RStudio functions.
- Environment Panel:** Shows the current environment is empty.
- Console:** Displays the R version 4.0.3 startup message and the R help text.

**File Explorer Window:**

- Path:** C:\Users\DChen\Desktop\DLR\RWorkshop
- Files:** .gitignore, README.md, RWorkshop.Rproj, 01\_IntroToR, 02\_DataWrangling, 03\_DataViz.

**Console Output:**

```
R version 4.0.3 (2020-10-10) -- "Bunny-wunnies Freak Out"
Copyright (C) 2020 The R Foundation for Statistical Computing
Platform: x86_64-w64-mingw32/x64 (64-bit)

R is free software and comes with ABSOLUTELY NO WARRANTY.
You are welcome to redistribute it under certain conditions.
Type 'license()' or 'licence()' for distribution details.

Natural language support but running in an English locale

R is a collaborative project with many contributors.
Type 'contributors()' for more information and
'citation()' on how to cite R or R packages in publications.

Type 'demo()' for some demos, 'help()' for on-line help, or
'help.start()' for an HTML browser interface to help.
Type 'q()' to quit R.

> |
```



# Running Code



## General Comments

If you see: `> print("Hello World")`

Run `print("Hello World")` in R Script. Do not include the ``>``.

Use ``#`` to write comments - code after `#` is not run.

```
> # This is not run
```

# Running Code in the Console

Type the following calculations:

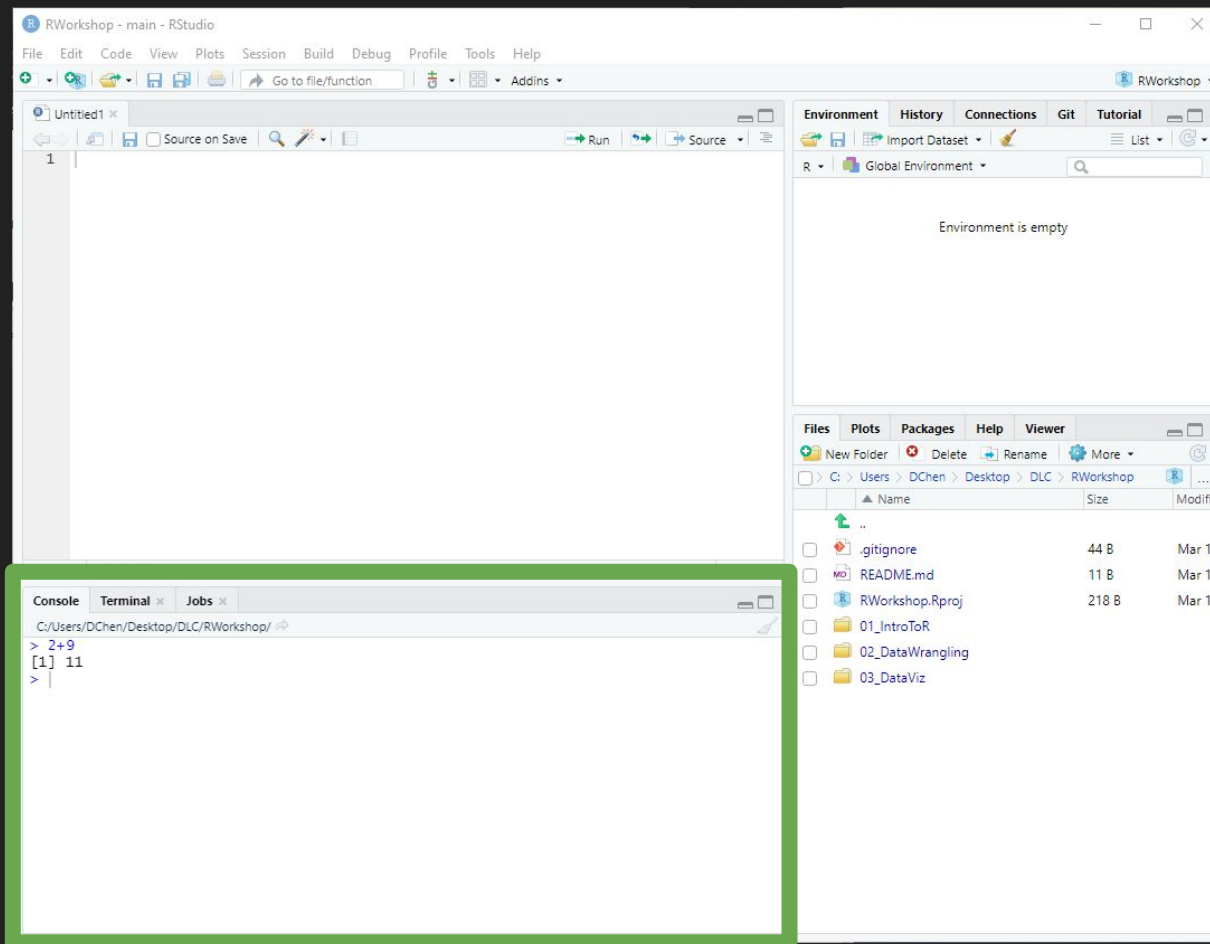
```
> 2+9
```

```
> 15*20
```

```
> 20/5
```

```
> 2^2
```

Press `Enter` to run



# Running Code in the Console

Type the following calculations:

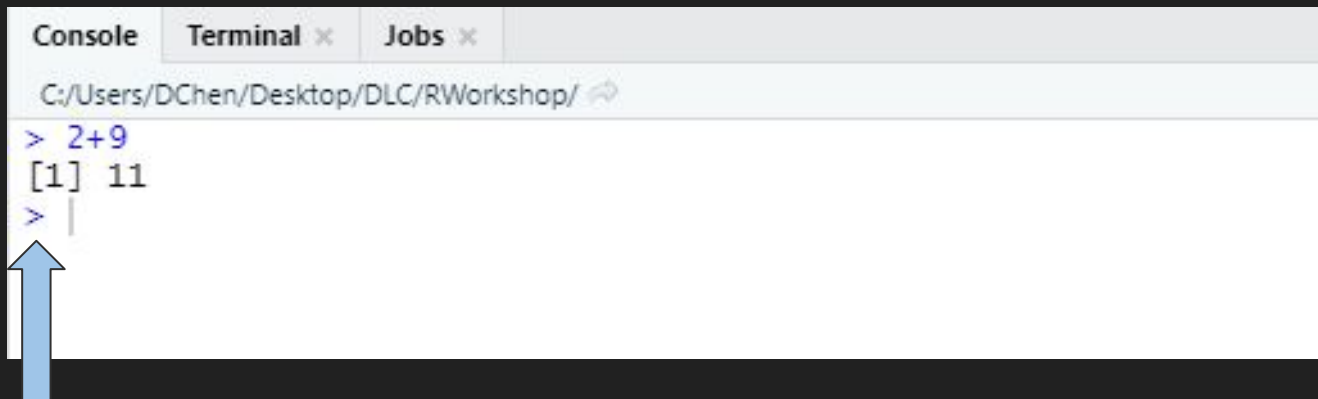
> 2+9

> 15\*20

> 20/5

> 2^2

Press `Enter` to run

A screenshot of a console window with tabs for 'Console', 'Terminal', and 'Jobs'. The console shows the R prompt '>' followed by the expression '2+9'. The output '[1] 11' is displayed on the next line. A blue arrow points to the prompt '>' on the third line, indicating where to enter new code.

```
Console Terminal x Jobs x
C:/Users/DChen/Desktop/DLC/RWorkshop/
> 2+9
[1] 11
> |
```

`>` means it's ready for new code

# Running Code in the Console

Type the following calculations:

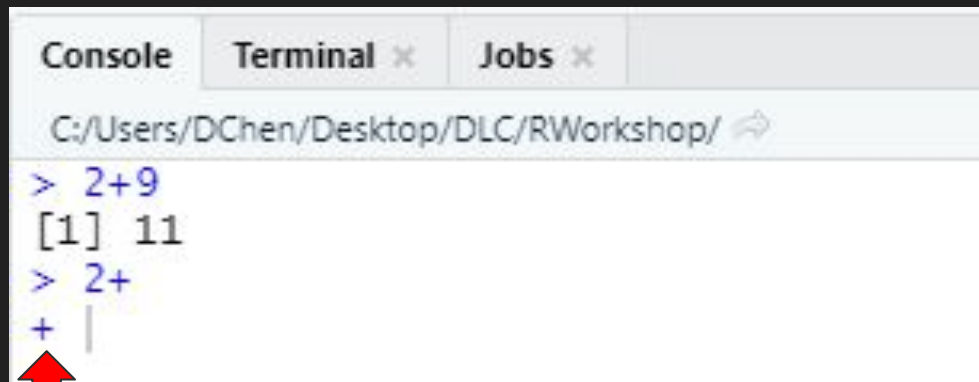
> 2+9

> 15\*20

> 20/5

> 2^2

Press `Enter` to run

A screenshot of a console window with tabs for 'Console', 'Terminal', and 'Jobs'. The 'Console' tab is active, showing the path 'C:/Users/DChen/Desktop/DLC/RWorkshop/'. The prompt '>' is followed by '2+9', and the output '[1] 11' is displayed. The prompt '>' is followed by '2+', and the output '+ |' is displayed. A red arrow points to the '+' character in the output.

```
Console Terminal x Jobs x
C:/Users/DChen/Desktop/DLC/RWorkshop/
> 2+9
[1] 11
> 2+
+ |
```

If you see + on the new line:

- Reset with the `Esc` key
- Continue from the previous code

# Running Code in the Console

Type the following calculations:

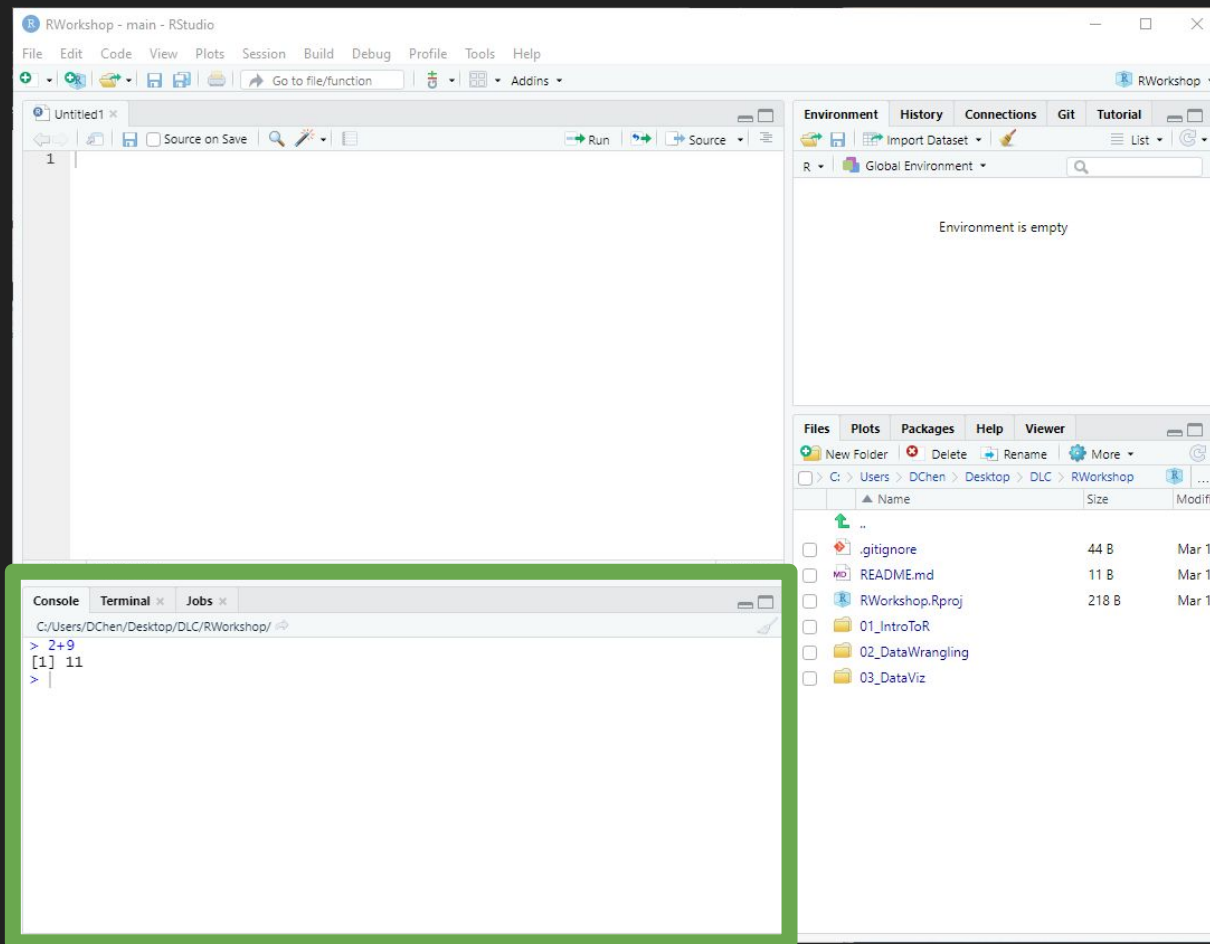
```
> 2+9
```

```
> 15*20
```

```
> 20/5
```

```
> 2^2
```

Press `Enter` to run



# Running Code in the Script

Code in the script can be saved and ran repeatedly

Output shows up in the console

The screenshot shows the RStudio interface. The top menu bar includes File, Edit, Code, View, Plots, Session, Build, Debug, Profile, Tools, and Help. Below the menu is a toolbar with icons for file operations and running code. The main editor window, titled 'Untitled1.R', contains the following R code:

```
1 2 + 9
2
3
4 15*20
5
6 20/5
7
8 2^2
9
```

The code is highlighted with a green border. The Environment pane on the right shows 'Global Environment' and 'Environment is empty'. The Files pane at the bottom right shows a directory structure with files like .gitignore, .Rhistory, 01\_IntroToR, 02\_DataWrangling, 03\_DataViz, README.md, and RWorkshop.Rproj. The Console pane at the bottom shows the R startup message and the output of the code:

```
C:/Users/DChen/Desktop/DLC/RWorkshop/
R is a collaborative project with many contributors.
Type 'contributors()' for more information and
'citation()' on how to cite R or R packages in publications.

Type 'demo()' for some demos, 'help()' for on-line help, or
'help.start()' for an HTML browser interface to help.
Type 'q()' to quit R.

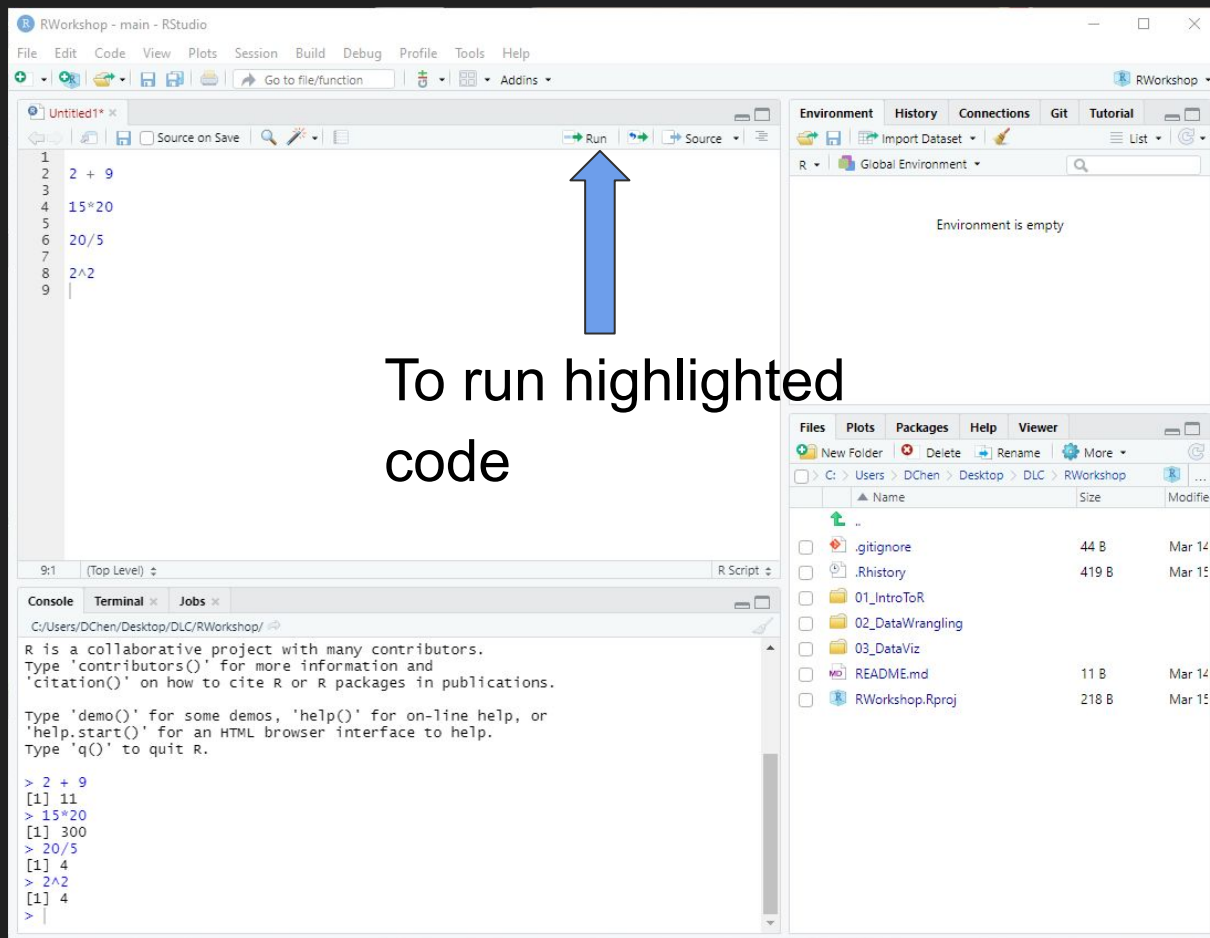
> 2 + 9
[1] 11
> 15*20
[1] 300
> 20/5
[1] 4
> 2^2
[1] 4
>
```

# Running Code in the Script

Type code into the R Script

Click or highlight the line of code and

- press `Run`
- press `CTRL` + `ENTER`



To run highlighted code

# Running Code in the Script

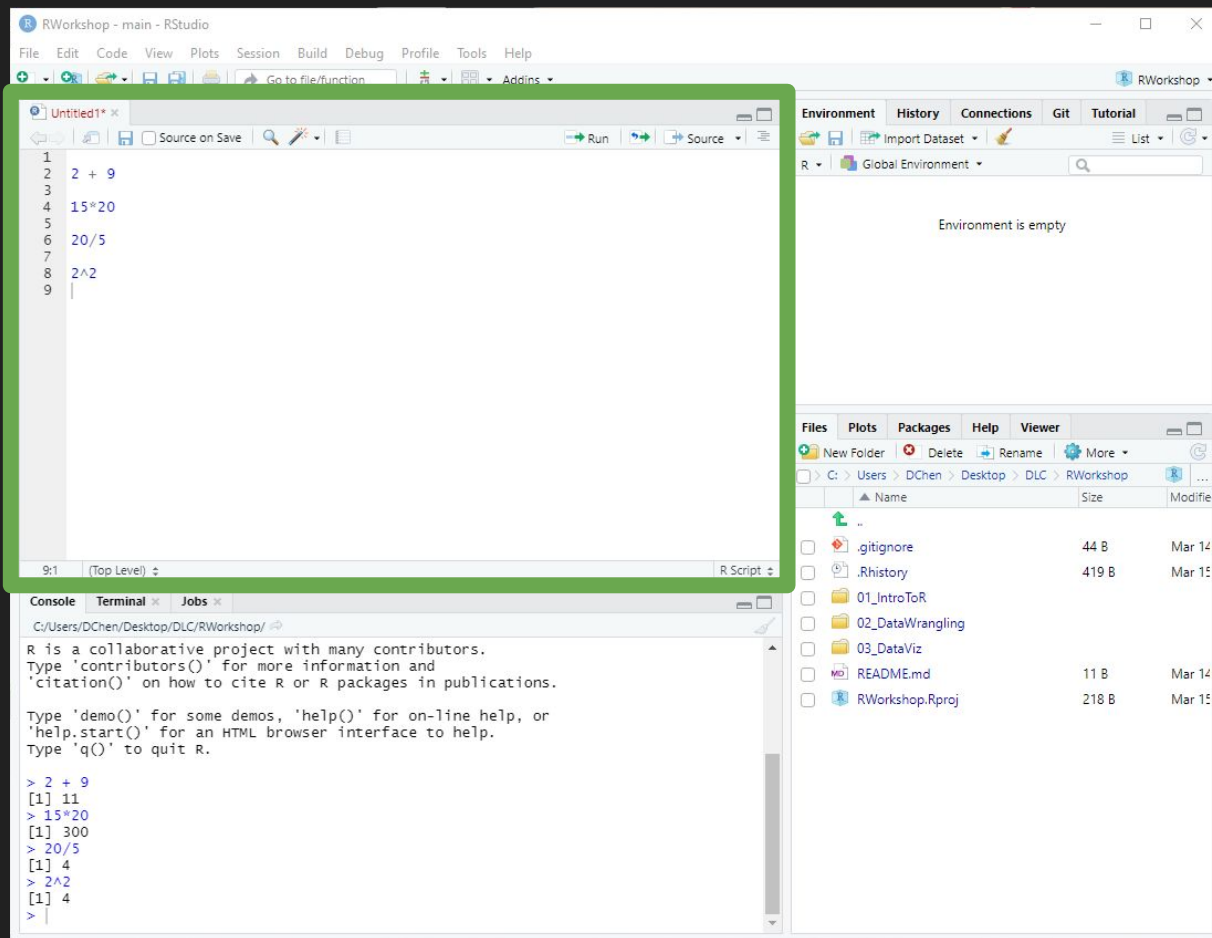
Type the following calculations into the R Script:

```
> 2+9
```

```
> 15*20
```

```
> 20/5
```

```
> 2^2
```





## Variable Assignment

variable <- value



Assignment  
Operator

(Can also be `=`)

## Variable Assignment

variable <- value

Replaces values

```
> x <- 5
```

```
> x + 10
```

## Variable Assignment

variable <- value

Replaces values

```
> x <- 5
```

```
> x + 10
```

Redefine over and  
over

```
> y <- 5
```

```
> y <- 3
```

```
> y
```

# Variable Assignment

variable <- value

Replaces values

```
> x <- 5
```

```
> x + 10
```

Redefine over and  
over

```
> y <- 5
```

```
> y <- 3
```

```
> y
```

Multiple  
variables

```
> x <- 1
```

```
> y <- 2
```

```
> x * y
```

## Variable Assignment - Comments

R is case sensitive

```
> value <- 5
```

```
> VALUE + 5
```

## Variable Assignment - Comments

R is case sensitive

```
> value <- 5
```

```
> VALUE + 5
```

Variable names must be  
one line and start with a  
character

```
> 1x <- 3      # not valid
```

```
> x y <- 3     # not valid
```

## Functions

```
function_name(input)
```

```
> sqrt(4)
```

```
> print("Hello World")
```

## Learn more about functions

- Google “<function name> in R”
- Use ``help()`` or ``?`` in R to see documentation

```
> help(sqrt)
```

```
> ?sqrt
```



# Data Types

## Common Data Types

Data Type	Examples
Numeric	-5, 1, 3.33, 100, pi

## Common Data Types

Data Type	Examples
Numeric	-5, 1, 3.33, 100, pi
Integer	-5L, 1L, 100L

## Common Data Types

Data Type	Examples
Numeric	-5, 1, 3.33, 100, pi
Integer	-5L, 1L, 100L
Character	'words', "3.33", 'TRUE', "1L"

## Common Data Types

Data Type	Examples
Numeric	-5, 1, 3.33, 100, pi
Integer	-5L, 1L, 100L
Character	'words', "3.33", 'TRUE', "1L"
Boolean/Logical	TRUE, FALSE, T, F

## Common Data Types

Data Type	Examples
Numeric	-5, 1, 3.33, 100, pi
Integer	-5L, 1L, 100L
Character	'words', "3.33", 'TRUE', "1L"
Boolean/Logical	TRUE, FALSE, T, F

## Common Data Types - Boolean

Operator	Description
<	Less than
<=	Less than or equal to
>	Greater than
>=	Greater than or equal to
==	Equal to
!=	Not equal to

## Common Data Types - Boolean

Examples	Operator	Description
TRUE	<	Less than
<ul style="list-style-type: none"><li>• 5 &gt; 3</li></ul>	<=	Less than or equal to
<ul style="list-style-type: none"><li>• 5 != 3</li></ul>	>	Greater than
FALSE	>=	Greater than or equal to
<ul style="list-style-type: none"><li>• 5 &lt;= 3</li></ul>	==	Equal to
<ul style="list-style-type: none"><li>• 8 == 10</li></ul>	!=	Not equal to



# Common Data Types - Boolean

Try it yourself!

```
> 5 > 3
```

```
> 10 == 10
```

Applies to variables!

```
> x <- 5
```

```
> x == 5
```

```
> x > 10
```

Operator

Description

<

Less than

<=

Less than or equal to

>

Greater than

>=

Greater than or  
equal to

==

Equal to

!=

Not equal to

## Common Data Types

Check the type of a variable/object with ``class()``:

```
> class("4")
```

```
> class(4)
```

Certain functions require specific data types:

```
> sqrt("4")
```

## Common Data Types - Common Mistakes

If a number is in quotations, think of it as the word instead of the number.

- “5” != 5

# Common Data Types - Common Mistakes

If a number is in quotations, think of it as the word instead of the number.

- `"5" != 5`

TRUE/FALSE are actually encoded as 1/0

- `TRUE == 1 ; FALSE == 0`

# Common Data Types - Common Mistakes

If a number is in quotations, think of it as the word instead of the number.

- `"5" != 5`

TRUE/FALSE are actually encoded as 1/0

- `TRUE == 1 ; FALSE == 0`

Characters without quotation marks are variables!

- `Hi <- 5`

# Common Data Types - Common Mistakes

If a number is in quotations, think of it as the word instead of the number.

- `"5" != 5`

TRUE/FALSE are actually encoded as 1/0

- `TRUE == 1 ; FALSE == 0`

Characters without quotation marks are variables!

- `Hi <- 5`

Be careful of single ``=``; outside of functions, it is assignment!

- `5 = 5 --> error!`
- `x = 5` is equivalent to `x <- 5`, but not recommended

# Data Structures

# Vectors

1	50	9	42
---	----	---	----

"A"	"B"	"C"
-----	-----	-----

TRUE	F
------	---

A column in excel

Any length from 1 onwards

All values have to be same type (all numeric, boolean, character, etc.)

Defined with function `c()`



# Vectors

1	50	9	42
---	----	---	----

"A"	"B"	"C"
-----	-----	-----

TRUE	F
------	---

A column in excel

Any length from 1 onwards

All values have to be same type (all numeric, boolean, character, etc.)

Defined with function `c()`

```
> lengths <- c(1, 50, 9, 42)
```

# Factors



A special vector with values assigned to each factor level (category)

## Data Frame

"A"	102	"Hela"	TRUE
"B"	40	"BHK"	F
"C"	12	"hESC"	T

Most common data format;  
equivalent to an excel sheet.

Multiple vectors combined  
together

Columns must be same type

All columns must have equal  
number of rows

# Reading in Data

Titanic Dataset: <https://web.stanford.edu/class/archive/cs/cs109/cs109.1166/stuff/titanic.csv>

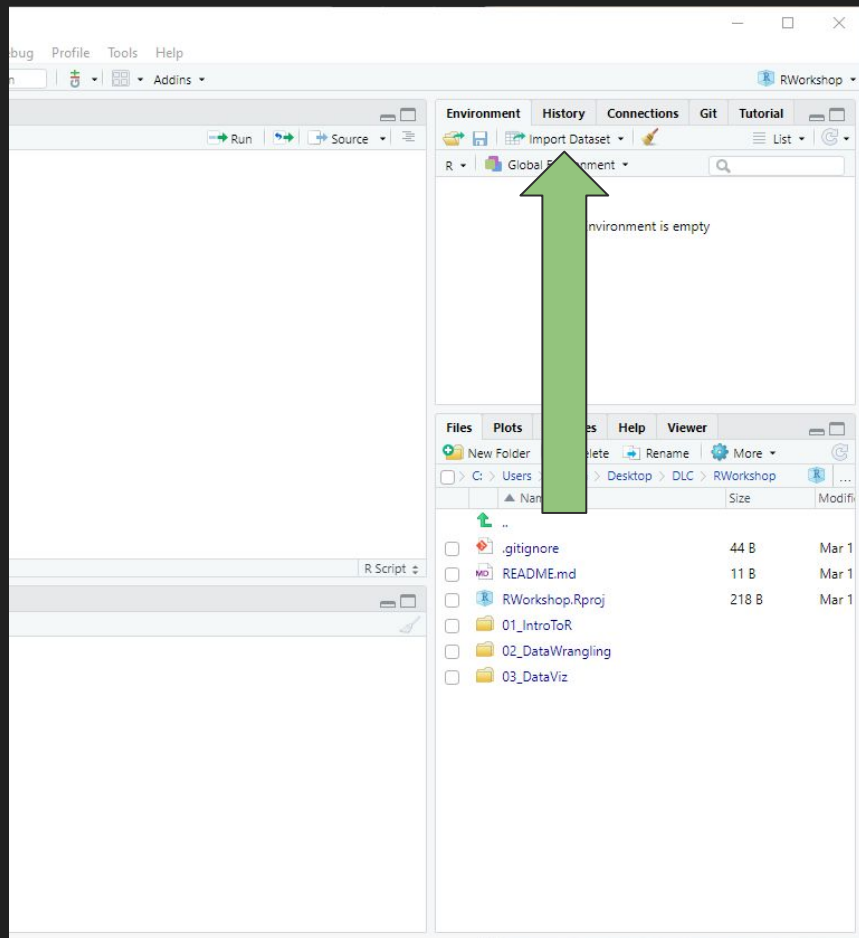
# Reading in Datasets

Data type	Extension	Function	Package
Comma separated values	csv	<code>read.csv()</code>	utils (default)
		<code>read_csv()</code>	readr (tidyverse)
Tab separated values	tsv	<code>read_tsv()</code>	readr
Other delimited formats	txt	<code>read.table()</code>	utils
Stata version 7-12	dta	<code>read.dta()</code>	foreign
SPSS	sav	<code>read.spss()</code>	foreign
SAS	sas7bdat	<code>read.sas7bdat()</code>	sas7bdat
Excel	xlsx, xls	<code>read_excel()</code>	readxl (tidyverse)

(HBC Source)

# Reading in Datasets

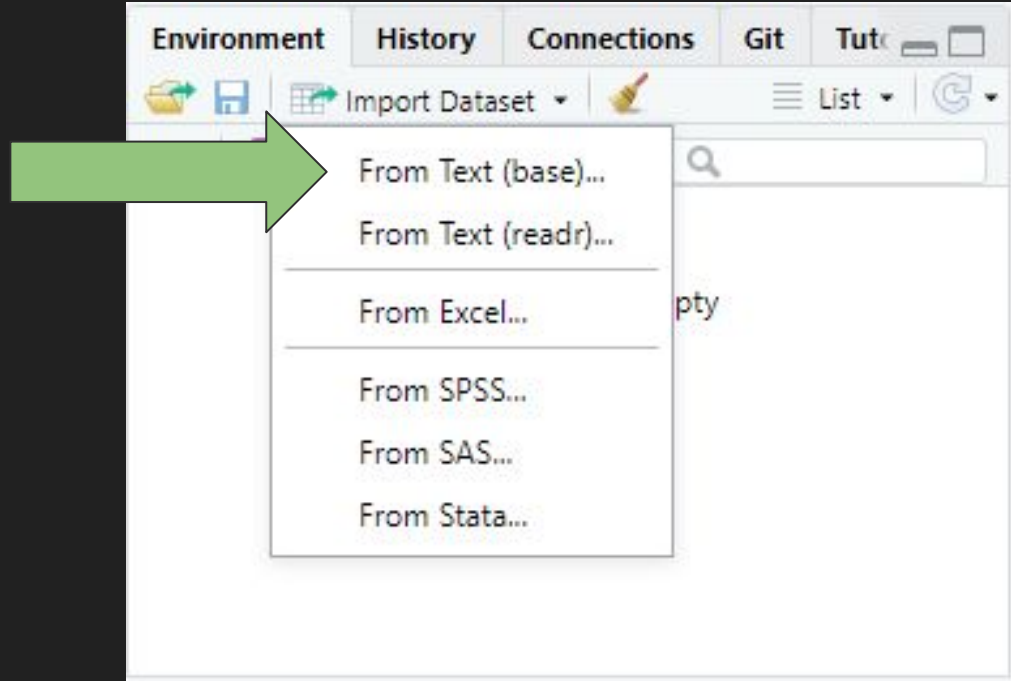
‘Import Dataset’ in the environment will automatically generate the code for you!



# Reading in Datasets

From Text works for almost all basic data files (.csv, .txt, etc.)

Find and select the dataset



# Reading in Datasets

Import Dataset

Name:

Encoding:

Heading: ☒ Yes ☐ No

Row names:

Separator:

Decimal:

Quote:

Comment:

na.strings:

☐ Strings as factors

Input File

Survived,Pclass,Name,Sex,Age,Siblings/Spouses Aboard,Parent

0,3,Mr. Owen Harris Braund,male,22,1,0,7.25

1,1,Mrs. John Bradley (Florence Briggs Thayer) Cumings,fem

1,3,Miss. Laina Heikkinen,female,26,0,0,7.925

1,1,Mrs. Jacques Heath (Lily May Peel) Futrelle,female,35,:

0,3,Mr. William Henry Allen,male,35,0,0,8.05

0,3,Mr. James Moran,male,27,0,0,8.4583

0,1,Mr. Timothy J McCarthy,male,54,0,0,51.8625

0,3,Master. Gosta Leonard Palsson,male,2,3,1,21.075

1,3,Mrs. Oscar W (Elisabeth Vilhelmina Berg) Johnson,femal

1,2,Mrs. Nicholas (Adele Achem) Nasser,female,14,1,0,30.07

1,3,Miss. Marguerite Rut Sandstrom,female,4,1,1,16.7

1,1,Miss. Elizabeth Bonnell,female,58,0,0,26.55

0,3,Mr. William Henry Saunderson,male,20,0,0,8.05

0,3,Mr. Anders Johan Andersson,male,39,1,5,31.275

0,3,Miss. Hulda Amanda Adolfina Vestrom,female,14,0,0,7.85

1,2,Mrs. (Mary D Kingcome) Hewlett,female,55,0,0,16

Data Frame

Survived	Pclass	Name
0	3	Mr. Owen Harris Braund
1	1	Mrs. John Bradley (Florence Briggs Thayer) Cumings
1	3	Miss. Laina Heikkinen
1	1	Mrs. Jacques Heath (Lily May Peel) Futrelle
0	3	Mr. William Henry Allen
0	3	Mr. James Moran
0	1	Mr. Timothy J McCarthy
0	3	Master. Gosta Leonard Palsson
1	3	Mrs. Oscar W (Elisabeth Vilhelmina Berg) Johnson
1	2	Mrs. Nicholas (Adele Achem) Nasser
1	3	Miss. Marguerite Rut Sandstrom
1	1	Miss. Elizabeth Bonnell
0	3	Mr. William Henry Saunderson
0	3	Mr. Anders Johan Andersson
0	3	Miss. Hulda Amanda Adolfina Vestrom
1	2	Mrs. (Mary D Kingcome) Hewlett

Import Cancel

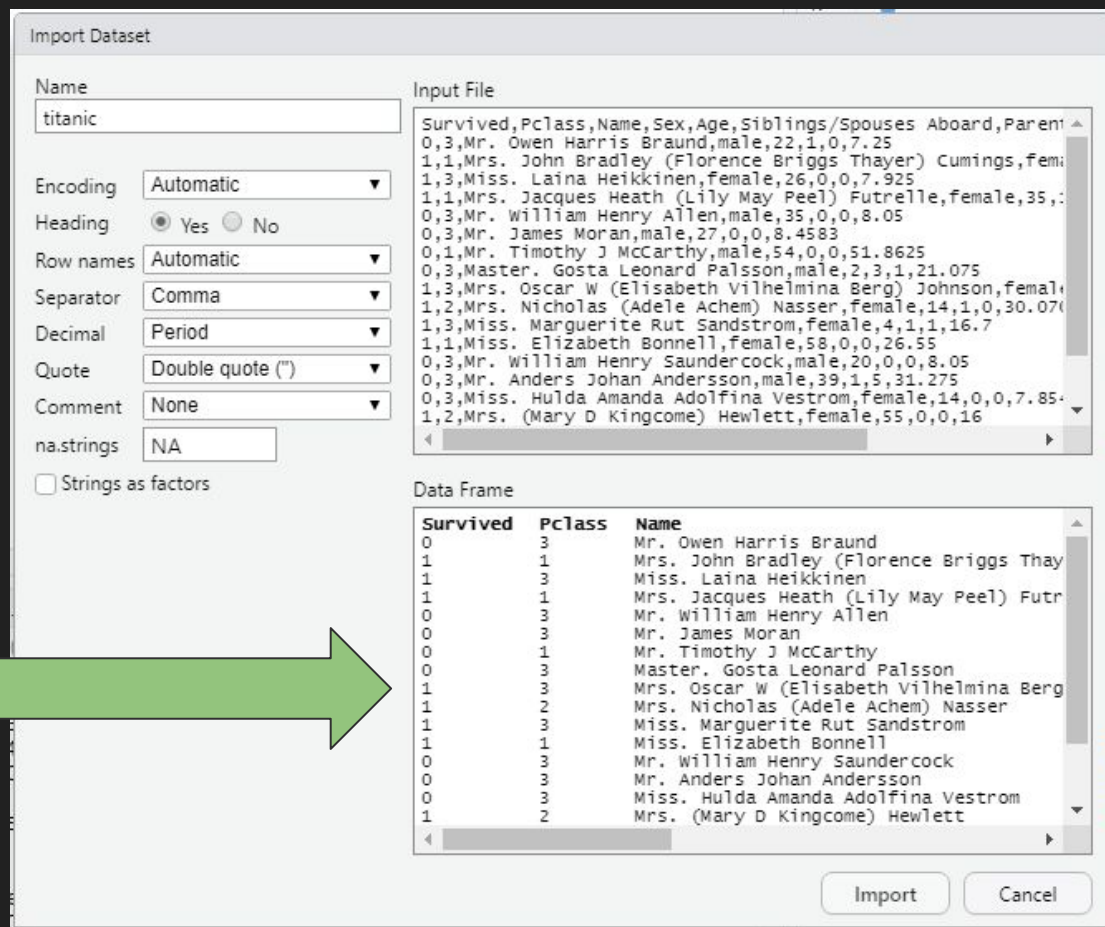


# Reading in Datasets

Make sure the data is being imported correctly!

Check:

- Headers
- Columns
- Values



Import Dataset

Name: titanic

Encoding: Automatic

Heading: ☒ Yes ☐ No

Row names: Automatic

Separator: Comma

Decimal: Period

Quote: Double quote (")

Comment: None

na.strings: NA

☐ Strings as factors

Input File: Survived,Pclass,Name,Sex,Age,Siblings/Spouses Aboard,Parent...

Data Frame:

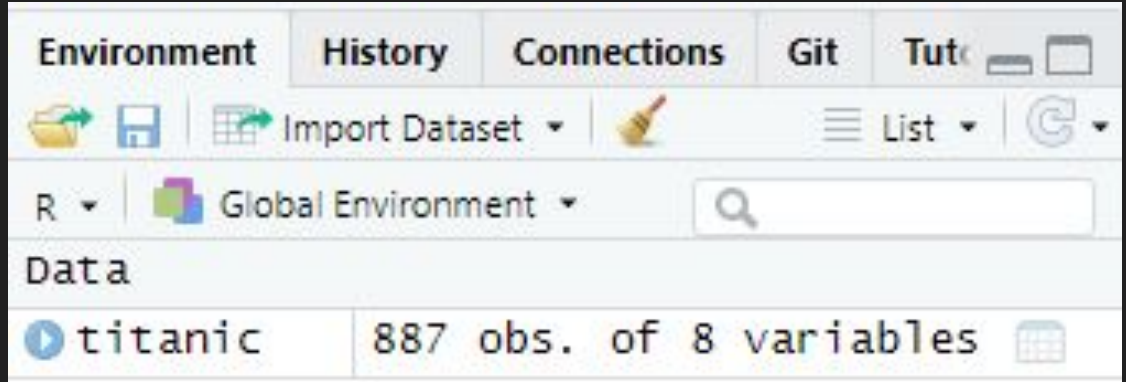
Survived	Pclass	Name
0	3	Mr. Owen Harris Braund
1	1	Mrs. John Bradley (Florence Briggs Thayer) Cumings,femi
1	3	Miss. Laina Heikkinen
1	1	Mrs. Jacques Heath (Lily May Peel) Futr
0	3	Mr. William Henry Allen
0	3	Mr. James Moran
0	1	Mr. Timothy J McCarthy
0	3	Master. Gosta Leonard Palsson
1	3	Mrs. Oscar W (Elisabeth Vilhelmina Berg) Johnson,femal
1	2	Mrs. Nicholas (Adele Achem) Nasser
1	3	Miss. Marguerite Rut Sandstrom
1	1	Miss. Elizabeth Bonnell
0	3	Mr. William Henry Saunderson
0	3	Mr. Anders Johan Andersson
0	3	Miss. Hulda Amanda Adolfin Vestrom
1	2	Mrs. (Mary D Kingcome) Hewlett

Import Cancel

## Viewing the Datasets

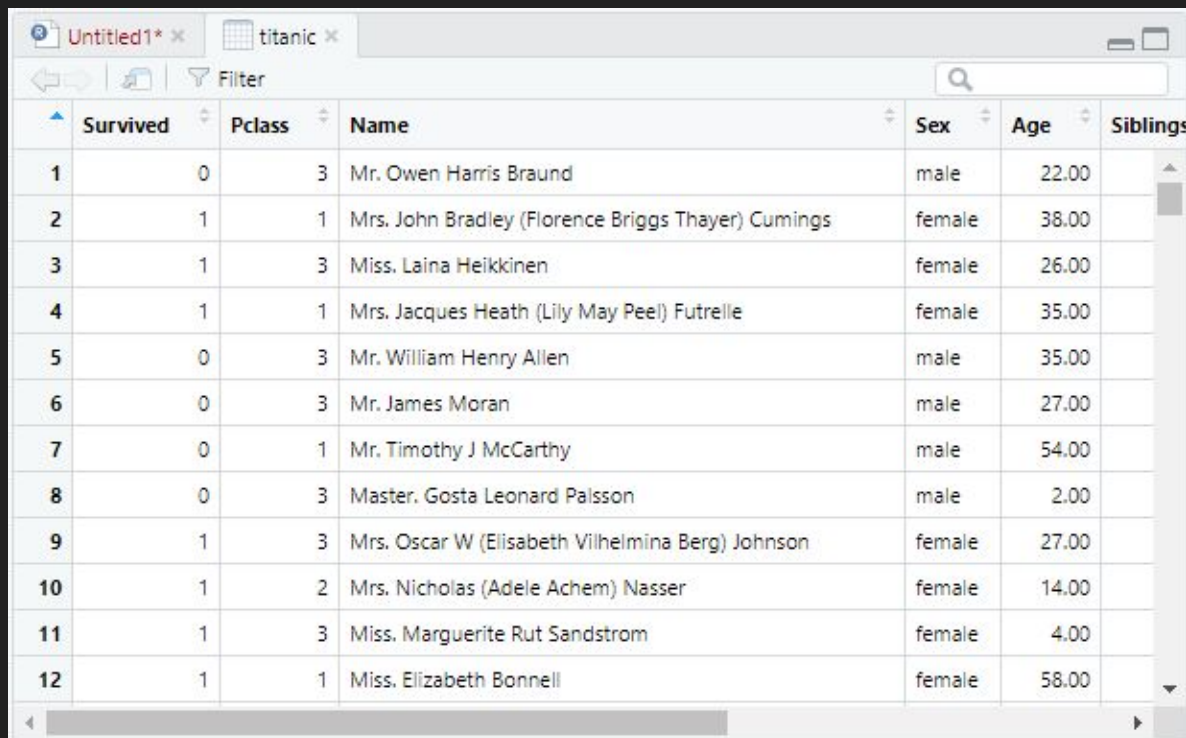
The data is now in the environment!

Click the name (titanic) to view the data!



# Viewing the Datasets

The data view after clicking the name in the environment



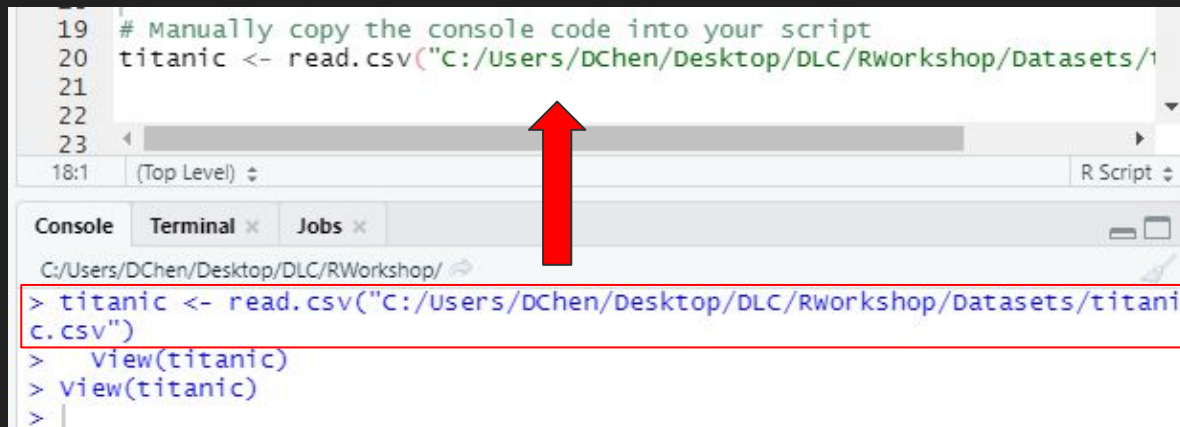
	Survived	Pclass	Name	Sex	Age	Siblings
1	0	3	Mr. Owen Harris Braund	male	22.00	
2	1	1	Mrs. John Bradley (Florence Briggs Thayer) Cumings	female	38.00	
3	1	3	Miss. Laina Heikkinen	female	26.00	
4	1	1	Mrs. Jacques Heath (Lily May Peel) Futrelle	female	35.00	
5	0	3	Mr. William Henry Allen	male	35.00	
6	0	3	Mr. James Moran	male	27.00	
7	0	1	Mr. Timothy J McCarthy	male	54.00	
8	0	3	Master. Gosta Leonard Palsson	male	2.00	
9	1	3	Mrs. Oscar W (Elisabeth Vilhelmina Berg) Johnson	female	27.00	
10	1	2	Mrs. Nicholas (Adele Achem) Nasser	female	14.00	
11	1	3	Miss. Marguerite Rut Sandstrom	female	4.00	
12	1	1	Miss. Elizabeth Bonnell	female	58.00	

## After Importing the Data

Copy the generated code into the R Script.

In the future, you will only need to run the code.

Do not include the `>` or the `View()` functions.



```
19 # Manually copy the console code into your script
20 titanic <- read.csv("C:/Users/DChen/Desktop/DLC/Rworkshop/Datasets/titanic.csv")
21
22
23
```

18:1 (Top Level) R Script

Console Terminal x Jobs x

C:/Users/DChen/Desktop/DLC/RWorkshop/

```
> titanic <- read.csv("C:/Users/DChen/Desktop/DLC/Rworkshop/Datasets/titanic.csv")
> view(titanic)
> view(titanic)
>
```

## Examine the Data - Try These Functions!

Functions - Add `titanic` to ()

---

str()

dim()

summary()

nrow()

head()

ncol()

tail()

colnames()

```
> str(titanic)
```

```
> head(titanic)
```

```
> colnames(titanic)
```

# Data Wrangling

## Before we start...

Import the `titanic` dataset into R

<https://web.stanford.edu/class/archive/cs/cs109/cs109.1166/problem12.html>

Install the packages ONCE (to check, skip this step)

```
> install.packages("tidyverse")
```

After installing, you must load the packages every time you open R

```
> library(tidyverse)
```

# What is the purpose of Data Wrangling?



# What is the Tidyverse?

Over 87 packages combined together

Packages align on data processing philosophy

Encompasses all aspects of a project, including data processing, visualization, modeling



<https://cran.r-project.org/web/packages/tidyverse/vignettes/paper.html>

# Tidyverse Philosophy

“Variables are in columns, observations are in rows, and values are in cells.”

country	year	cases	population
Afghanistan	1999	745	19987071
Afghanistan	2000	8666	20593360
Brazil	1999	37737	172006362
Brazil	2000	80488	174604898
China	1999	212258	1272915272
China	2000	213766	128042583

variables

country	year	cases	population
Afghanistan	1999	745	19987071
Afghanistan	2000	8666	20593360
Brazil	1999	37737	172006362
Brazil	2000	80488	174604898
China	1999	212258	1272915272
China	2000	213766	128042583

observations

country	year	cases	population
Afghanistan	1999	745	19987071
Afghanistan	2000	8666	20593360
Brazil	1999	37737	172006362
Brazil	2000	80488	174604898
China	1999	212258	1272915272
China	2000	213766	128042583

values

# The Fundamental Functions



# The Fundamental Functions

`select`

Select columns

## select - Select which variables to keep

```
select(data, `Column1`, `Column2`,...)
```



Columns to keep or  
remove (-)

Use column name,  
number, or colon

## select - Keep columns

`select(data, Name)`

Name	Major	HW
Matt	math	60
Mary	math	90
Bill	biology	85
Brie	biology	87

Select column  
`Name`



Name
Matt
Mary
Bill
Brie

## select - Remove Columns

NOT

`select(data, -Major)`

Name	Major	HW
Matt	math	60
Mary	math	90
Bill	biology	85
Brie	biology	87

Select all columns  
**except** `Major`

Name	HW
Matt	60
Mary	90
Bill	85
Brie	87

## select - Keep Multiple Columns

```
select(data, Major, HW)
```

Name	Major	HW
Matt	math	60
Mary	math	90
Bill	biology	85
Brie	biology	87

Select columns  
`Major` and `HW`



Major	HW
math	60
math	90
biology	85
biology	87



## select - Keep Multiple Columns

```
select(data, Major:HW)
```

Name	Major	HW
Matt	math	60
Mary	math	90
Bill	biology	85
Brie	biology	87

Select columns from  
`Major` to `HW`



Major	HW
math	60
math	90
biology	85
biology	87

## select - Titanic Exercises

1. Select the variables `Survived`, `Sex`, and `Parents.Children.Aboard`

```
> select(titanic, _____, _____, _____)
```

2. Remove the `Name` variable, and store the new results in `Titanic\_nameless`

```
> Titanic_nameless <- select(titanic, _____)
```

3. Store the first four columns of `titanic` in a new dataframe. How many different ways can you do this using `select()`?

## select - Titanic Exercise Solutions

1. Select the variables `Survived`, `Sex`, and `Parents.Children.Aboard`

```
> select(titanic, Survived, Sex, Parents.Children.Aboard)
```

2. Remove the `Name` variable, and store the new results in `Titanic\_nameless`

```
> Titanic_nameless <- select(titanic, -Name)
```

3. Store the first four columns of `titanic` in a new dataframe. How many different ways can you do this using `select()`?

## select - Titanic Exercise Solutions

3. Store the first four columns of `titanic` in a new dataframe. How many different ways can you do this using `select()`?

```
> t2 <- select(titanic, 1:4)
```

```
> t2 <- select(titanic, Survived:Sex)
```

```
> t2 <- select(titanic, Survived, Pclass, Name, Sex)
```

```
> t2 <- select(titanic, -Age, -Siblings.Aboard, -Parents.Children.Aboard, -Fare)
```

```
> t2 <- select(titanic, -c(Age:Fare))
```



# The Fundamental Functions

`select`

Select columns



# The Fundamental Functions

select

Select columns

filter

Filter rows

## filter - filter the rows

`filter(data, <Boolean condition>, ...)`

Condition on the  
variables

## filter - filter the rows

filter(data, <Boolean condition>,...)

Condition on the  
variables

<	Less than	>=	Greater than or equal to
<=	Less than or equal to	==	Equal to
>	Greater than	!=	Not equal to



## filter - Keep specific value

```
filter(data, Major == "math")
```

Name	Major	HW
Matt	math	60
Mary	math	90
Bill	biology	85
Brie	biology	87

Keep rows where  
Major is "math"



Name	Major	HW
Matt	math	60
Mary	math	90

## filter - Keep range of values

```
filter(data, HW <= 85)
```

Name	Major	HW
Matt	math	60
Mary	math	90
Bill	biology	85
Brie	biology	87

Keep rows where `HW`  
is less than or equal to  
85



Name	Major	HW
Matt	math	60
Bill	biology	85

## filter - Multiple Conditions

```
filter(data, HW <= 85, Major == "math")
```

Name	Major	HW
Matt	math	60
Mary	math	90
Bill	biology	85
Brie	biology	87

Keep rows where `HW` is  
less than or equal to 85,  
**AND** `Major` is math



Name	Major	HW
Matt	math	60

## filter - Multiple Conditions

AND

```
filter(data, HW <= 85 & Major == "math")
```

Name	Major	HW
Matt	math	60
Mary	math	90
Bill	biology	85
Brie	biology	87

Keep rows where `HW` is  
less than or equal to 85,  
**AND** `Major` is math

Name	Major	HW
Matt	math	60

## filter - Multiple Conditions

OR

```
filter(data, HW <= 85 | Major == "biology")
```

Name	Major	HW
Matt	math	60
Mary	math	90
Bill	biology	85
Brie	biology	87

Keep rows where `HW` is  
less than or equal to 85,  
**OR** `Major` is biology

Name	Major	HW
Matt	math	60
Bill	biology	85
Brie	biology	87

## filter - Titanic Exercises

1. Create a new dataset with only those who survived (`Survived` value of 1)

```
> survived <- filter(titanic, _____ == _____)
```

2. Only show rows with 1st class (`Pclass`) females (`Sex`)

```
> filter(titanic, _____ == 1 _ Sex == _____)
```

3. How many rows are there if you keep only males between ages 20-30?

## filter - Titanic Exercise Solutions

1. Create a new dataset with only those who survived (`Survived` value of 1)

```
> survived <- filter(titanic, Survived == 1)
```

```
> survived <- filter(titanic, Survived == "1")
```

## filter - Titanic Exercise Solutions

1. Create a new dataset with only those who survived (`Survived` value of 1)

```
> survived <- filter(titanic, Survived == 1)
```

```
> survived <- filter(titanic, Survived == "1")
```

2. Only show rows with 1st class (`Pclass`) females (`Sex`)

```
> filter(titanic, Pclass == 1 & Sex == "female")
```

```
> filter(titanic, Pclass == 1, Sex == "female")
```



## filter - Titanic Exercise Solutions

3. How many rows are there if you keep only males between ages 20-30?

```
> tm <- filter(titanic, Sex == "male", Age >= 20, Age <= 30)
```

```
> nrow(tm) # 223
```

```
> t2 <- filter(titanic, Sex == "male")
```

```
> t3 <- filter(t2, Age >= 20)
```

```
> t4 <- filter(t3, Age <= 30)
```



# The Fundamental Functions

select                      Select columns

filter                      Filter rows

# Pipes

%>%

## Pipes %>%

Pipes (%>%) insert the data into the next line

---

function(data)		data %>%
		function(.)

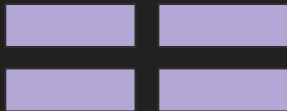
## Pipes %>%

sqrt(2)



2 %>% sqrt()

sqrt(sqrt(2))



2 %>%

sqrt() %>%

sqrt()

## Pipes %>%

Without pipes, multi-line processing is hard to read

```
df1 <- select(df, A, B, C)
df2 <- filter(df1, A == 1, B > 2)
df3 <- filter(df2, C == 1 | C == 5)
```

Creates unnecessary  
intermediate datasets

## Pipes %>%

Without pipes, multi-line processing is hard to read

```
df1 <- select(df, A, B, C)
```

```
df2 <- filter(df1, A == 1, B > 2)
```

```
df3 <- filter(df2, C == 1 | C == 5)
```

```
df1 <- filter(filter(select(df, A, B, C), A ==  
1, B > 2), C == 1 | C == 5)
```

Hard to read



## Pipes %>%

With pipes, one line is one logical step

---

```
df1 <- df %>%
```

Start with dataset `df`



## Pipes %>%

With pipes, one line is one logical step

---

```
df1 <- df %>%
```

Start with dataset `df`

```
  select(A, B, C) %>%
```

Select variables A, B, C

## Pipes %>%

With pipes, one line is one logical step

---

```
df1 <- df %>%
```

Start with dataset `df`

```
  select(A, B, C) %>%
```

Select variables A, B, C

```
  filter(A == 1, B > 2) %>%
```

Keep rows with A == 1 and B > 2

## Pipes %>%

With pipes, one line is one logical step

---

```
df1 <- df %>%
```

Start with dataset `df`

```
  select(A, B, C) %>%
```

Select variables A, B, C

```
  filter(A == 1, B > 2) %>%
```

Keep rows with A == 1 and B > 2

```
  filter(C == 1 | C == 5)
```

Keep rows with C == 1 or C == 5

(Store results in df1)

## Pipes %>%

df: 

A	B	C	D	E
---	---	---	---	---

%>%

select(A, B, C) %>%

A	B	C
---	---	---

select(A)

%>%

new\_df:

A
---

## Pipe %>% - Titanic Exercise

1. Using pipes, store the first four columns and keep only 1st class (`Pclass`)

```
> firstclass <- _____ %>%
```

```
> select(1_4) _____
```

```
> filter(Pclass _____)
```

## Pipe %>% - Titanic Exercise Solution

1. Using pipes, store the first four columns and keep only 1st class (`Pclass`)

```
> firstclass <- titanic %>%
```

```
>   select(1:4) %>%
```

```
>   filter(Pclass == 1)
```



# The Fundamental Functions

select

Select columns

filter

Filter rows



## The Fundamental Functions

`select`                      Select columns

`filter`                      Filter rows

`mutate`                      Create new variables



## **mutate - Create new variables/columns**

```
mutate(., new_variable = value, ...)
```

Computation of new variable  
Can be based on another variable,  
a function, a single value, etc.

## mutate - Create var with constant value

```
data %>% mutate(Sect = 1)
```

Name	HW	Ex
Matt	60	75
Mary	90	100
Bill	85	60
Brie	87	85

Create new variable  
`Sect` where all values  
are 1



Name	HW	Ex	Sect
Matt	60	75	1
Mary	90	100	1
Bill	85	60	1
Brie	87	85	1

## mutate - Create var based on another var

```
data %>% mutate(Ex2 = Ex/100)
```

Name	HW	Ex
Matt	60	75
Mary	90	100
Bill	85	60
Brie	87	85

Create new variable `Ex2`  
as `Ex` divided by 100



Name	HW	Ex	Ex2
Matt	60	75	0.75
Mary	90	100	1.00
Bill	85	60	0.60
Brie	87	85	0.85

## mutate - Create var based on other variables

```
data %>% mutate(Final = (HW + Ex)/2)
```

Name	HW	Ex
Matt	60	75
Mary	90	100
Bill	85	60
Brie	87	85

Create new variable  
`Final` as average of  
`HW` and `Ex`



Name	HW	Ex	Final
Matt	60	75	67.5
Mary	90	100	95
Bill	85	60	72.5
Brie	87	85	86

## mutate - Titanic Exercises

1. `Fare` is currently in pounds. Convert it to dollars (1 pound = 1.37 dollars)

```
> titanic %>% mutate(FareDollars = _____ * _____)
```

2. How many family members total (combine `Siblings.Spouses.Aboard` and `Parents.Children.Aboard`) were on board for each passenger?

```
> titanic %>% mutate(titanic, Family = _____ + _____)
```

3. Using pipes: remove the name, keep only 1st class survivors, subtract 20 from `Age`. Store these results in a dataset named `AgeAdjusted`.

## mutate - Titanic Exercise Solutions

1. `Fare` is currently in pounds. Convert it to dollars (1 pound = 1.37 dollars)

```
> titanic %>% mutate(FareDollars = Fare*1.37)
```

2. How many family members total (combine `Siblings.Spouses.Aboard` and `Parents.Children.Aboard`) were on board for each passenger?

```
> titanic %>% mutate(Family =  
Siblings.Spouses.Aboard + Parents.Children.Aboard)
```

## mutate - Titanic Exercise Solutions

3. Using pipes: remove the name, keep only 1st class survivors, subtract 20 from `Age`. Store these results in a dataset named `AgeAdjusted`.

```
> AgeAdjusted <- titanic %>%
```

```
>   select(-Name) %>%
```

```
>   filter(Pclass == 1, Survived == 1) %>%
```

```
>   mutate(Age = Age - 20)
```



## The Fundamental Functions

`select`                      Select columns

`filter`                      Filter rows

`mutate`                      Create new variables





## The Fundamental Functions

`select`                      Select columns

`filter`                      Filter rows

`mutate`                      Create new variables

`summarize`                  Compute summary variable

## summarize - Create new aggregated variables

```
summarize(., new_variable = function(_), ...)
```

Computation of new variable  
Function that returns one value  
(e.g. mean(), median(), max())



**summarize - Create new aggregated variables**

summarize() computes a single value or column

mutate() adds a variable to existing dataset

## summarize - Compute mean of dataset

```
data %>% summarize(HW_mean = mean(HW))
```

Name	HW
Matt	60
Mary	90
Bill	85
Brie	87

Calculate mean HW  
score



HW_mean
80.5

## summarize - Count total number of observations

```
data %>% summarize(Count = n())
```

Name	HW
Matt	60
Mary	90
Bill	85
Brie	87

Calculate number of  
observations



Count
4



## The Fundamental Functions

`select`                      Select columns

`filter`                      Filter rows

`mutate`                      Create new variables

`summarize`                  Compute summary variable



## The Fundamental Functions

<code>select</code>	Select columns
<code>filter</code>	Filter rows
<code>mutate</code>	Create new variables
<code>summarize</code>	Compute summary variable
<code>group_by</code>	Split data by groups

## group\_by - split data by groups

```
group_by(., `id_var`, ...)
```

Split dataset by `id\_var` (or multiple variables)

Dataset split based on unique values of `id\_var`



## group\_by - split data by groups

```
data %>% group_by(., Major)
```

Major	HW
math	60
math	90
biology	85
biology	87

## group\_by - split data by groups

```
data %>% group_by(., Major)
```

Major	HW
math	60
math	90
biology	85
biology	87

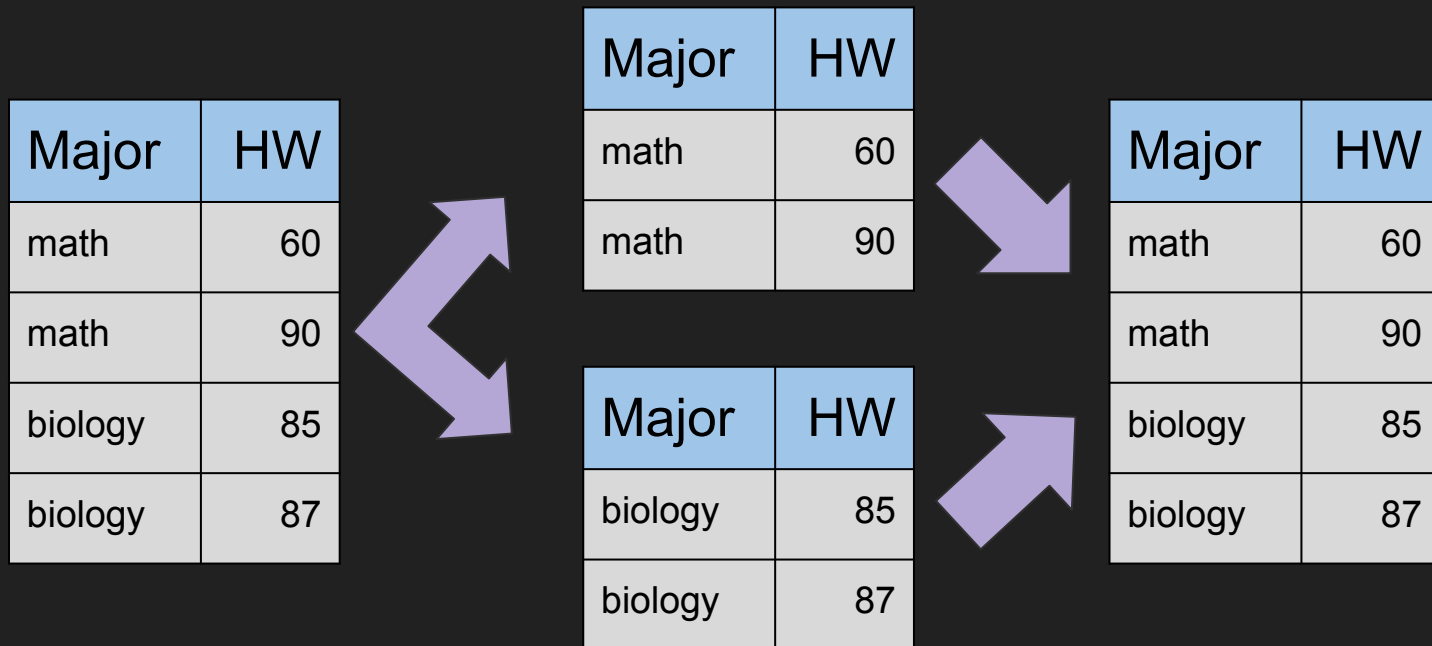


Major	HW
math	60
math	90

Major	HW
biology	85
biology	87

## group\_by - split data by groups

```
data %>% group_by(., Major)
```



## group\_by() - split data by groups

group\_by(., Major)

Major	HW
math	60
math	90
biology	85
biology	87

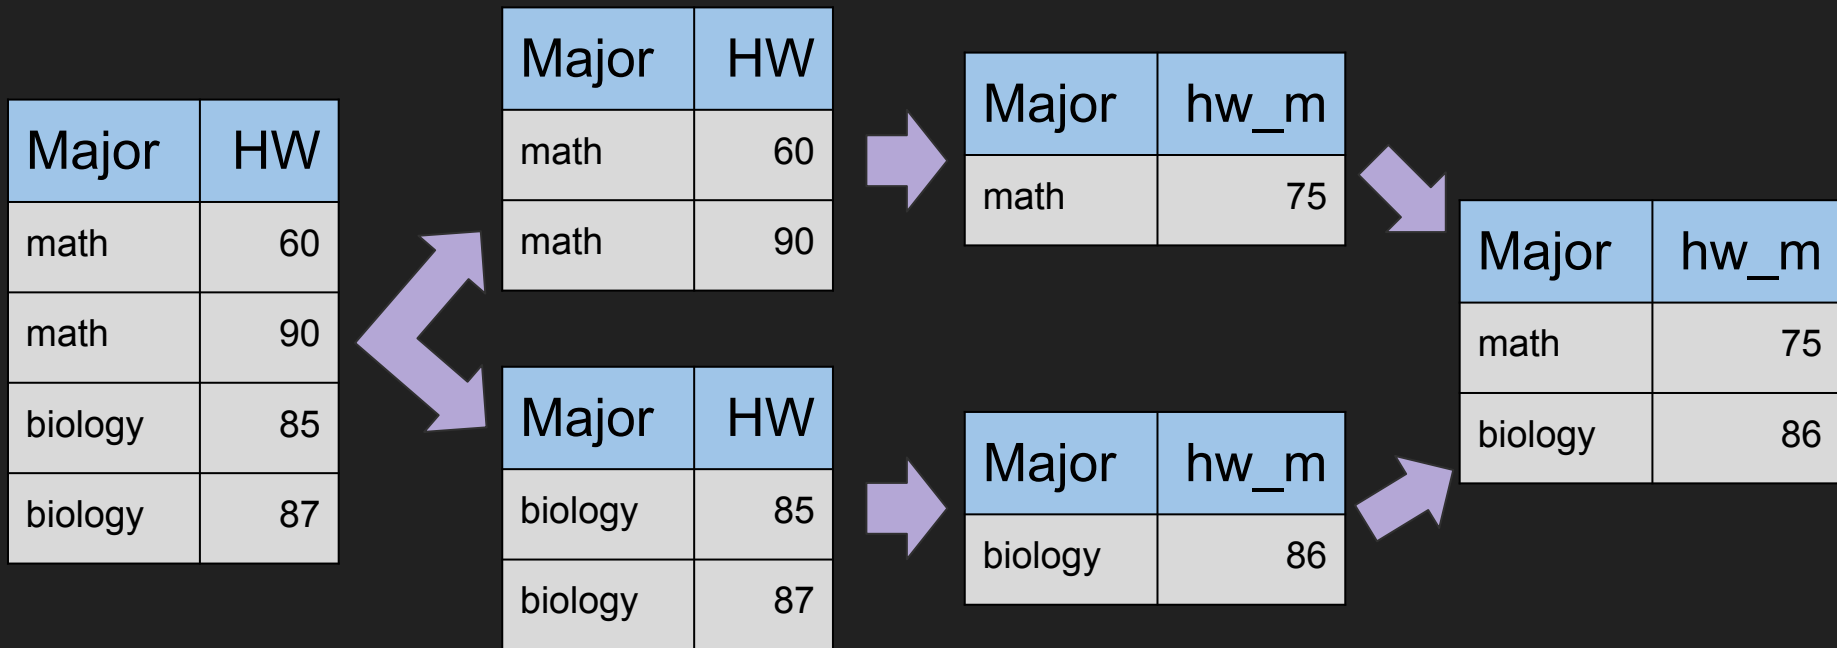


Major	HW
math	60
math	90

Major	HW
biology	85
biology	87

## group\_by() - split data by groups

```
group_by(., Major) %>% summarize(hw_m = mean(HW))
```



## group\_by() - split data by groups

group\_by(., Major)

Major	HW
math	60
math	90
biology	85
biology	87



Major	HW
math	60
math	90

Major	HW
biology	85
biology	87

## group\_by() - split data by groups

Take top HW value

```
group_by(., Major) %>% top_n(1, HW)
```

Major	HW
math	60
math	90
biology	85
biology	87



Major	HW
math	60
math	90



Major	HW
math	90



Major	HW
math	90
biology	87

Major	HW
biology	85
biology	87



Major	HW
biology	87



## summarize() and group\_by() - Titanic Exercises

1. What was the average `Fare` cost?

```
> titanic %>% summarize(FareMean = mean(_____))
```

2. What was the average `Age` for each `Pclass`? Store the results in `ClassAge`.

```
> ClassAge <- titanic %>%
```

```
> group_by(_____) %>% summarize(_____ = _____(_____))
```

3. What percent of each `Pclass` survived? (Hint: sum() and n() are helpful).



## summarize() and group\_by() - Titanic Exercise Solutions

1. What was the average `Fare` cost?

```
> titanic %>% summarize(FareMean = mean(Fare))
```

## summarize() and group\_by() - Titanic Exercise Solutions

1. What was the average `Fare` cost?

```
> titanic %>% summarize(FareMean = mean(Fare))
```

2. What was the average `Age` for each `Pclass`? Store the results in `ClassAge`.

```
> ClassAge <- titanic %>%
```

```
> group_by(Pclass) %>%
```

```
> summarize(AgeM = mean(Age))
```

## summarize() and group\_by() - Titanic Exercises

3. What percent of each `Pclass` survived? (Hint: sum() and n() are helpful).

```
> titanic %>%
```

```
>   group_by(Pclass) %>%
```

```
>   summarize(pct = sum(Survived)/n())
```



## The Fundamental Functions

<code>select</code>	Select columns
<code>filter</code>	Filter rows
<code>mutate</code>	Create new variables
<code>summarize</code>	Compute summary variable
<code>group_by</code>	Split data by groups



## The Fundamental Functions

<code>select</code>	Select columns
<code>filter</code>	Filter rows
<code>mutate</code>	Create new variables
<code>summarize</code>	Compute summary variable
<code>group_by</code>	Split data by groups
<code>left_join</code>	Combine datasets

## `left_join()` - combine two datasets

```
left_join(df_left, df_right, ...)
```

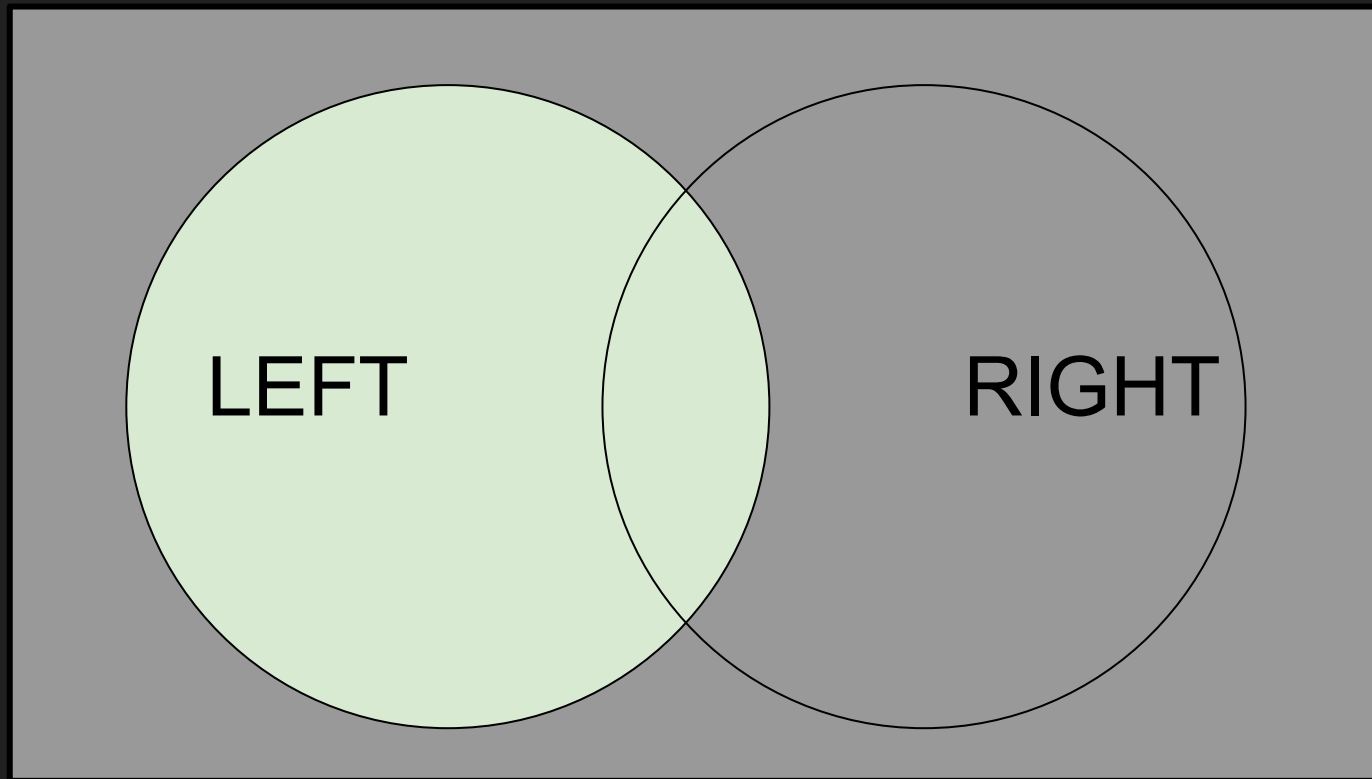
The two datasets to combine  
together

Join two tables, matching on a key column

## `left_join()` - combine two datasets

All values in the  
left dataset are  
kept

Only matching  
values in the  
right dataset are  
kept



## left\_join() - combine two datasets

```
Student <- left_join(Majors, Grades)
```

Majors

Name	Major
Matt	math
Mary	math
Bill	biology
Brie	biology



Grades

Name	HW
Matt	60
Mary	90
Bill	85
Brie	87



Student

Name	Major	HW
Matt	math	60
Mary	math	90
Bill	biology	85
Brie	biology	87



## left\_join() - combine two datasets

```
Student <- left_join(Majors, Grades)
```

Majors

Name	Major
Matt	math
Mary	math
Bill	biology
Brie	biology



Grades

Name	HW
Matt	60
Mary	90



Student

Name	Major	HW
Matt	math	60
Mary	math	90
Bill	biology	NA
Brie	biology	NA

## left\_join() - combine two datasets

```
Student <- left_join(Majors, Grades)
```

Majors

Name	Major
Matt	math
Mary	math



Grades

Name	HW
Matt	60
Mary	90
Bill	85
Brie	87



Student

Name	Major	HW
Matt	math	60
Mary	math	90

## left\_join() - combine two datasets

If matching columns  
have different names

```
Student <- left_join(Grades, Majors, by = c("NX" = "NY"))
```

Majors

NX	Major
Matt	math
Mary	math
Bill	biology
Brie	biology



Grades

NY	HW
Matt	60
Mary	90
Bill	85
Brie	87



Student

NX	Major	HW
Matt	math	60
Mary	math	90
Bill	biology	85
Brie	biology	87

## Other joins

Function	Description
<code>left_join</code>	Keep all left
<code>right_join</code>	Keep all right
<code>inner_join</code>	Match left and right
<code>full_join</code>	Keep all values
<code>nest_join</code>	Create list with match
<code>anti_join</code>	Keep non-matching

## left\_join() - Titanic Exercise

1. Split the titanic dataset accordingly:

```
> titanic1 <- titanic %>% select(1:3)
```

```
> titanic2 <- titanic %>% select(3:5)
```

2. Combine the datasets back together based on `Name`

```
> titanic3 <- left_join(____, ____)
```

## left\_join() - Titanic Exercise

1. Split the titanic dataset accordingly:

```
> titanic1 <- titanic %>% select(1:3)
```

```
> titanic2 <- titanic %>% select(3:5)
```

2. Combine the datasets back together based on `Name`

```
> titanic3 <- left_join(titanic1, titanic2)
```

## Exporting Data

Name in R



```
write_csv(dataset, file = "dataset.csv")
```



Output name and path to  
location

**NOTE: MUST INCLUDE .CSV**

## Exporting Data - Titanic Exercises

1. Store the first four columns in `TitanicReduced`

```
> TitanicReduced <- titanic %>% select(_____)
```

2. Export the previously created dataset

```
> write_csv(_____, "RTitanic.csv")
```



## Exporting Data - Titanic Exercise Answers

1. Store the first four columns in `TitanicReduced`

```
> TitanicReduced <- titanic %>% select(1:4)
```

2. Export the previously created dataset

```
> write_csv(TitanicReduced, "RTitanic.csv")
```



## Other Tips

Working with dates/times? Use the ``lubridate`` package!

Skim through the data wrangling cheat sheet to see what is possible.

Do not reinvent the wheel. Use google before you try to code something fancy.

**Have fun!**